

Sebastian Thrun

- planning, perception, control

- Team: Anthony Novarro
Andy Brown
Ceromne Camacho
Elica White
Taran Ziggee
Jessica Bulovics
Erika Alonso

David Silver
↳

- map

- Localization

- plan → choose a trajectory amongst all possibilities.

- shortest path
(split) Decision

- search problem

- Darpa Grand Challenge

- Red Whittakers

Sebastian Thrun

Mike Montemerlo

Anthony Levandowski

$$a_{\text{lat}} = \frac{V^2}{R}$$

lateral acceleration:

In the Junkyard conditions the maximum lateral acceleration that Corki's tires support is 12 m/s^2 . If she travels at a velocity of 30 m/s what is the sharpest turn (minimum turning radius) she can make?

75 m

129, 30 → stop

125, 40 → full gas back, full steer right

till 126, 5, 34
full gas back, full steer left

- noise (not deterministic, every run results in a slightly different behavior)

- no accuracy down to desired level.
(you say stop at position 40, it actually stops at 38)

- difficult to test

- the less parameters, the better
(e.g. ~~considering n and y~~ at the same time makes it more difficult)

- update rate of information is very important

- windows of action ranges of values with probability.

Bayes Rule:

- Normal robots needs to know everything about the world. When the world is wrong, then it would do fatal mistakes.
- Bayes Rule permits for uncertainty; that the robot might not know and might have to explore the world in order to get knowledge about it.

- conditional probability
- Bayes Rule
- Probability distribution (+ Bayes rule) \Rightarrow
- robot localizer \hookrightarrow is crucial for robot localization
- \hookrightarrow Proj 1: 2D Histogram Filter.

- Uncertainty / control noise

Uncertainty

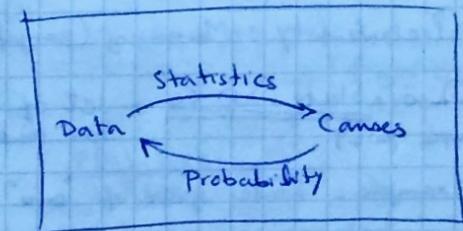
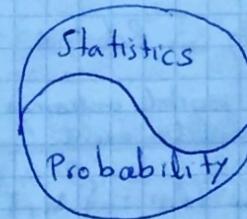
- "knowable" in probabilistic sense \neq certain
- uncertain \rightarrow **Nothing is certain!** in a SDC

- \hookrightarrow what traffic around the car does in near future
- location of the car
- speed of the car
- Steering
- ...

- complete certainty vs. sufficient certainty

- A car is an imperfect mechanical system

- Our brain is a high performance adaptive controller



Probability: Gives us the language for describe, discuss the relationship between Data \leftrightarrow Causes data and the underlying causes.

Let's "make data" by flipping the coin:

Data = {H, H, H} \rightsquigarrow Data from three flips (measurement)

Probability as a method of describing the anticipated outcome of these coin flips.

- Fair Coin vs. Biased Coin

$$P(A) = 1 - P(\neg A)$$

- A coin flip is an example of a "Probabilistic Event"

- Probabilistic Event: a set of outcomes to some experiment where each outcome has a probability.

- Example: Coin Flipping

\hookrightarrow outcomes: H, T
 $\hookrightarrow P(H) = 0.5, P(T) = 0.5$

- Truth Table: listing e.g. all possible outcomes of a 2 coin flip

Flip 1 | Flip 2
H | H
H | T
T | H
T | T

Probability: Managing Complexity in Robotics

- Ω • State Space: set of all possible outcomes for (sample space) for a probabilistic event.

Example: State space of "one" coin flip: {H, T}

- Calculation ~~of sets~~ needed for N events each having a space state of size n is n^N

→ Exponential complexity growth

- Probability $\approx \frac{\text{Number of ways an event can occur}}{\text{Total number of events that could occur}}$

Conditional Probability:

the basic mechanism that a SDC can take some sensor data and infer something else. For example it could see a picture of a tree and infer "I am about to crash".

The language that is needed in mathematics to infer from sensor measurements to some internal belief is called conditional probability.

- Frequent sensor measurement → more "Fresh" data
→ better estimate of conditional probability
 $P(\text{Event} | \text{measurement})$

- "Cloudiness now" and "cloudiness one minute from now" are not independent events.

- We take sensor measurements to use that data to inform our probabilistic estimate of an event.

(Intuition) Conditional Probability:

Taking advantage of what we know, to make better estimates about what we don't.

$$P(A|B) = \frac{P(A, B)}{P(B)}$$

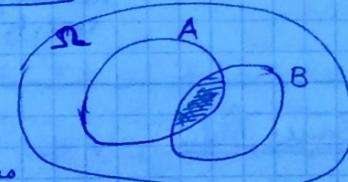
If A and B independent:

$$P(A|B) = \frac{P(A, B)}{P(B)} = \frac{P(A) \cdot P(B)}{P(B)} \Rightarrow$$

$P(A|B) = P(A)$ if A and B are independent

$$P(A|B) = \frac{P(A, B)}{P(B)}$$

Intuition:



Intuition: Restricting the sample space to situations in which ~~A or B~~ B occurs; Because if possible outcomes for A and B are restricted to those in which B occurs, this set serves as the new sample space

$$P(A|B) = \frac{P(A, B)}{P(B)}$$

$$P(B|A) = \frac{P(A, B)}{P(A)}$$

$$\Rightarrow P(A|B) \cdot P(B) = P(B|A) \cdot P(A)$$

$$\Rightarrow P(B|A) = \frac{P(A|B) \cdot P(B)}{P(A)}$$

C: Cancer

Cancer and Test example from lecture: Test: P or N

$$P(C) = 0.9$$

$$P(\neg C) = 0.1$$

$$P(P|C) = 0.9$$

$$P(N|C) = 0.1$$

$$P(P|\neg C) = 0.2$$

$$P(N|\neg C) = 0.8$$

Total Probability: ← Total Probability

$$P(P) = P(P|C) \cdot P(C) + P(P|\neg C) \cdot P(\neg C)$$

	Cancer	Test	Truth Table
C	P	→ P(C C) = 0.09	
¬C	P	→ P(N C) = 0.01	
C	¬P	→ P(C \neg C) = 0.18	
¬C	¬P	→ P(N \neg C) = 0.82	
			Σ = 1

Law of Total Probability:

Example: Coin $\sim T, H$: $P(T) = P(H) = \frac{1}{2}$

1. flip $\rightarrow T$ & don't flip
2. flip $\rightarrow H$? \rightarrow flip again.

what is the prob. of H at second flip?

$$P(H) = P(H|T)P(T) + P(H|H')P(H')$$

$$= 0 \cdot \frac{1}{2} + \frac{1}{2} \cdot \frac{1}{2} = \frac{1}{4}$$

If B_1, B_2, B_3, \dots is a partition of a sample space S , then for any event A :

$$P(A) = \sum_i P(A \cap B_i) = \sum_i P(A|B_i) * P(B_i)$$

(think of it as: calculating the sum of all probabilities necessary to ensure all scenarios for a specific event are included.)

Bayes Rule

Reducing Uncertainty

$$P(A|B) = \frac{P(B|A) \cdot P(A)}{P(B)}$$

$$P(C_k|x) = \frac{P(C_k) \cdot P(x|C_k)}{P(x)}$$

posterior \propto Prior \cdot likelihood
evidence

• ~~Robotics~~ SLDs make a lot of measurements for example about speed, location of the car that are not certain. These measurements are dependent on one another

• Bayes rule gives us a tool to combine the (on each other depending information) from different sensors and combine them into a more reliable representation of the car's position, movement, and environment.

using conditional probabilities

Bayes Rule

Reducing Uncertainty

- Bayes rule is very important in robotics.

- Bayes rule summarized in one sentence:

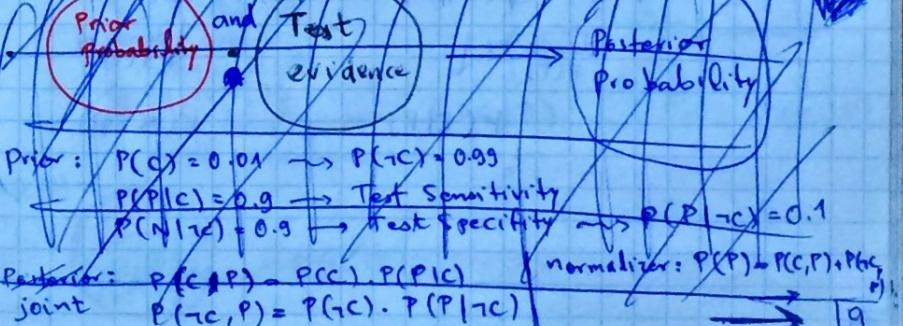
"Given an initial prediction, if we gather additional related data, data that our initial prediction ~~depends~~ depends on, we can improve that prediction"

- initial prediction \approx prior belief

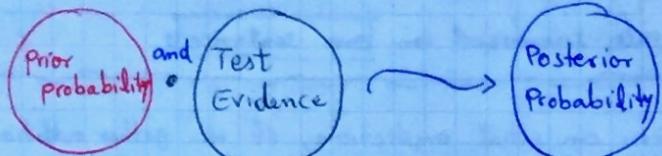
- The sensor data combined with what we already know about the road and the car gives us more information about where our location is most likely to be
- Bayes rule gives us a mathematical way to correct our measurements and let's us move from an uncertain prior belief to something more and more probable.

Cancer Test: Sensitivity $\approx P(P|C)$
Specificity $\approx P(N|\neg C)$

Essence of Bayes Rule:



Essence of Bayes rule:



Prior: $P(C) = 0.01 \rightarrow P(\neg C) = 0.99$
 $P(P|C) = 0.9$
 $P(N|\neg C) = 0.9 \rightarrow P(P|\neg C) = 0.1$

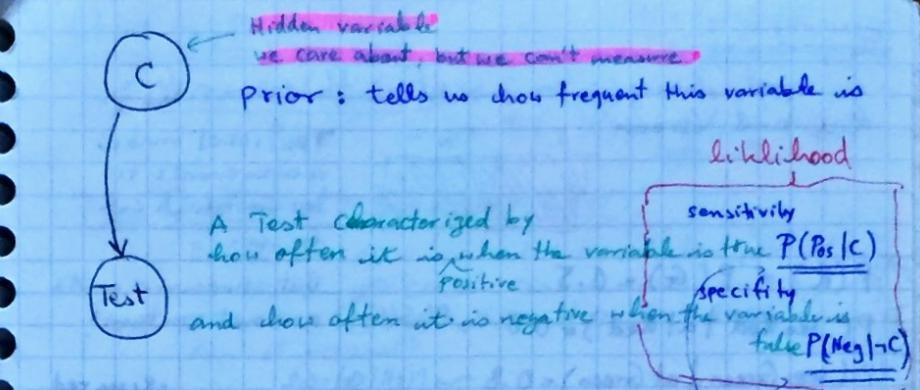
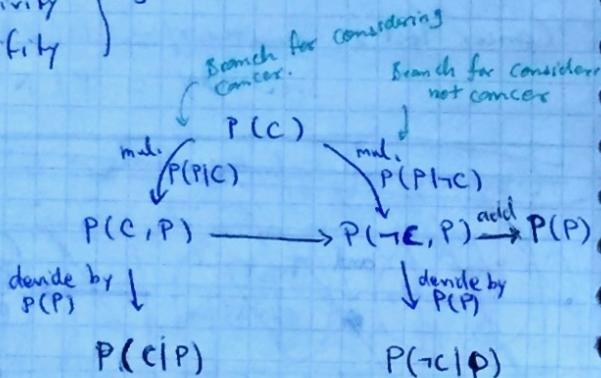
Joint: $P(C, P) = P(C) \cdot P(P|C) = 0.009$
 $P(\neg C, P) = P(\neg C) \cdot P(P|\neg C) = 0.099$

Normalizer: $P(P) = P(C, P) + P(\neg C, P) = 0.108$

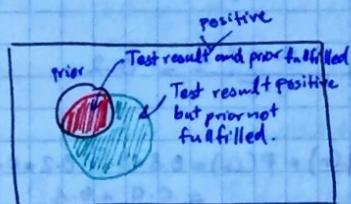
Posterior: $P(C|P) = \frac{P(C, P)}{P(P)} = 0.0833$ $P(\neg C|P) = \frac{P(\neg C, P)}{P(P)} = 0.9167$

$P(C)$ prior
 $P(P|C)$ test sensitivity
 $P(N|\neg C)$ test specificity

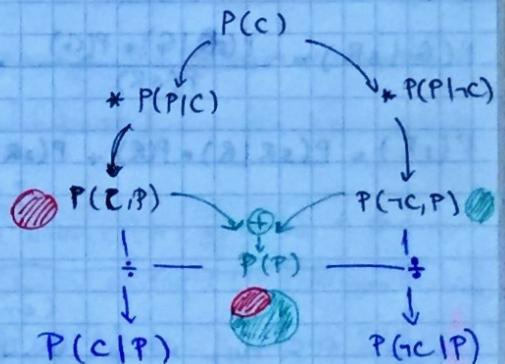
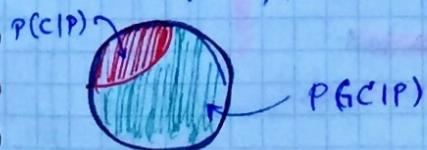
Received: P



Bayes rule now takes the prior, multiplies in the measurement frequencies to give us a new variable $P(C, P)$ and the same for measurement given the opposite assumption about the hidden variable, giving us $P(\neg C, P)$, we add those to come to $P(P)$ normalizer and we divide $P(C, P)$ and $P(\neg C, P)$ by normalizer to arrive at our best estimate of the hidden variable C given our test result.

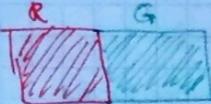


But the red and green areas don't add up to 1! therefore divide by total area



Example:

A robot living in a world with two cells:
Red and Green



The robot makes measurements to see in which cell he is.

$$P(R) = P(G) = 0.5$$

$P(\text{see Red} | \text{at Red}) = 0.8 \rightsquigarrow P(sG|R)=0.2$ are not certain!

$$P(\text{see Green} | \text{at Green}) = 0.8 \rightsquigarrow P(sR|G)=0.2$$
 sR: see red

Assume Robot sees red.

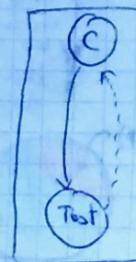
- What is now the posterior probability that the robot is at red cell? $0.8 \quad P(R|sR)$
- What is the probability that it is in green cell even though it sees red? $0.2 \quad P(G|sR)$

$$P(R|sR) = \frac{P(sR|R) * P(R)}{P(sR)} = \frac{0.8 * 0.5}{0.5} = 0.8$$

$$P(G|sR) = \frac{P(sR|G) * P(G)}{P(sR)} = \frac{0.2 * 0.5}{0.5} = 0.2$$

$$P(sR) = P(sR|R) * P(R) + P(sR|G) * P(G) = 0.8 * 0.5 + 0.2 * 0.5 = 0.4 + 0.1 = 0.5$$

Consider
We are now able to ~~test~~ let a
hidden variable
and understand how a test
can give us information back
about it!



Steps to Bayes Rule:

H = Hypothesis
T = Test

- Prior Probabilities: Based on previous data, how likely is that H happens?
 $P(H) \rightsquigarrow P(\neg H)$ likely is that H doesn't happen?

2. Conditional / Test Probabilities:

Based on sensor or test data collection, how likely is that a ~~certain~~ certain test or sensor reading is to occur, given that the hypothesis H does or has not occurred?

$$\begin{aligned} P(T|H) &\rightsquigarrow P(\neg T|H) \\ P(T|\neg H) &\rightsquigarrow P(\neg T|\neg H) \end{aligned}$$

3. Joint Probabilities: of the prior and the test:

$$P(T, H) = P(T|H) \cdot P(H)$$

$$\text{likewise: } P(T, \neg H) = P(T|\neg H) \cdot P(\neg H)$$

4. Total Probabilities:

- Total probability of a test result or sensor reading.
- This is the sum of joint probabilities in which that test result occurs.
- We need this to normalize the posterior probability.

$$P(T) = P(T, H) + P(T, \neg H)$$

5. Posterior Probability:

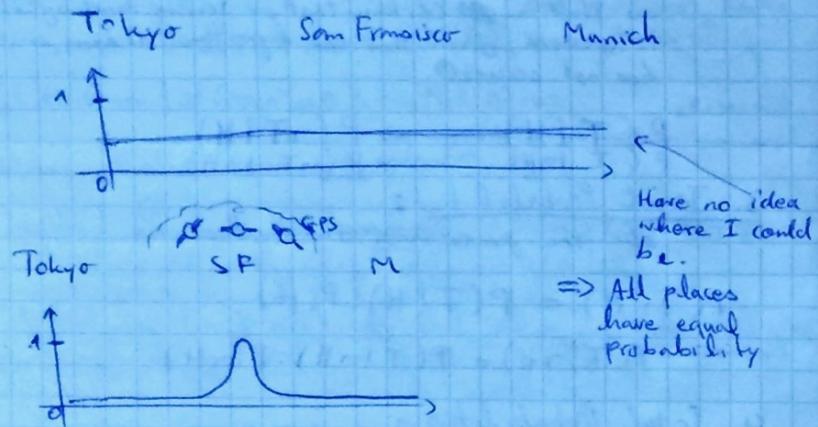
Probability of an event (H) given the ~~a~~ certain test or a certain sensor reading (T) [Bayes Rule]

$$P(H|T) = \frac{P(T|H) \cdot P(H)}{P(T)} \quad P(H|T) = \frac{P(T, H)}{P(T)}$$

posterior \propto likelihood · prior evidence

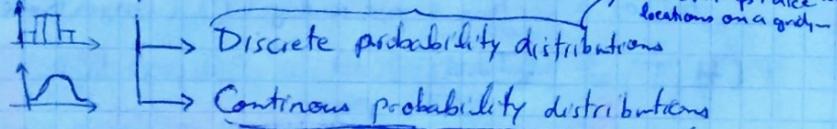
Probability Distributions

- A way in which a SDC represents its internal beliefs about the world.
- A mathematical way to represent uncertainty across all possible outcomes



- Very useful in visualizing probabilities.
- Can be worked with using algebra, calculus, linear algebra
- Probability distribution allows us to show the probability of an event using a mathematical formula.

Two Types:



- I Discrete prob. dist. &
- the variable can take discrete values
 - all values y of p are greater or equal to 0
 - for each y , $p(y) = y$
 - sum of all y values is 1

Continuous Prob. dist.

! Every outcome (exact outcome) has prob. 0 !!

$$P(\text{angle} = 180) = 0$$

$$P(\text{angle} = 179.999) = 0$$

Every outcome has $p = 0$!

Because it is impossible to have the outcome be exactly equal the given value!

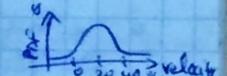


X: outcome

$$P(0 \leq x \leq 180) = 0.5$$

$$P(260 \leq x \leq 290) = \frac{1}{12} = 0.08333$$

$$P(179 \leq x \leq 180) = \frac{1}{360} = 0.00277$$



Probability Density Function PDF

- Characteristics of continuous prob. dist. and Prob. density func
 - y values are greater or equal zero
 - prob. of specific x occurring is equal zero
 - prob. of an event occurring between two values of x , x_1 and x_2 is equal the area under the curve between x_1 and x_2
 - The total area under the prob. dens. func. is 1
- For a continuous dist. we can only calculate the probs. between a range of values !

Uniform Continuous Dist. :

- PDF \rightarrow an equation that mathematically represents a continuous distribution.

$$(P = \int \text{PDF})$$

Densities can be larger than 1

But the sum under the continuous curve must add up to 1

```
import matplotlib.pyplot as plt
```

```
x = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]  
y = [2, 4, 6, 8, 10, 12, 14, 16, 18, 20]
```

```
plt.scatter(x,y) // plt.bar(x,y) // plt.pdot(x,y) //  
plt.xlabel("x values") // plt.plot([0,360],  
plt.ylabel("y values") // [y,y],  
plt.title("X values versus Y values") // "—",  
plt.xticks([0,1,2,3,4,5,6,7,8,9,10,11]) // color="r")  
plt.show()
```

```
// plt.imshow(grid,  
            cmap="F310",  
            clim=(0,1))
```

Gaussian Distributions

- There are many different continuous prob. dists.
- For each application, we use a suitable dist
- We use Gaussian dist. to model uncertainty in sensor measurements.



PDF for Gaussian dist:

$$f(n) = \frac{1}{\sqrt{2\pi}\sigma^2} e^{-\frac{(n-\mu)^2}{2\sigma^2}}$$

μ : mean of population

σ : standard dev. of population

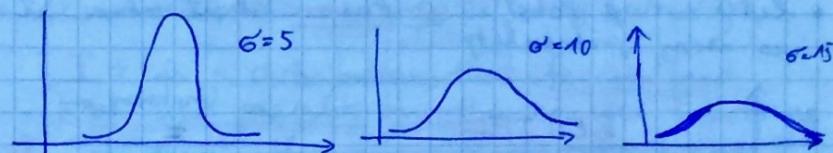
$$f(n) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{n-\mu}{\sigma}\right)^2}$$



- μ : shifts the curve along the x axis.

- σ : makes the curve taller or wider.

\Rightarrow As the standard deviation increases, the curve becomes wider and flatter.
 \Rightarrow the uncertainty increases



Central limit Theorem:

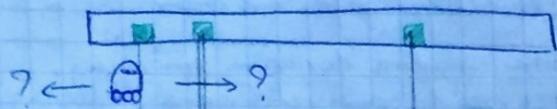
- Says: \rightarrow for independent, random variables If you get large enough samples from a population and then calculate the sample's mean, these means will be normally distributed

Robot Localization:

- GPS \rightarrow 2-10 m accuracy
- For a SDC \rightarrow we need 2-10 cm
- The techniques here require a good map. If we don't have a map, we have to do SLAM.

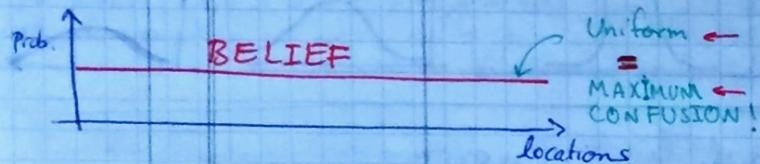


Imagine a 1D world



■ = Door

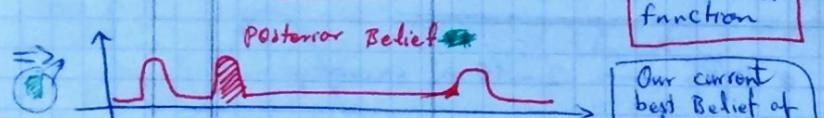
- let's map robot's confusion about where it is using probability



- to be able to localize we have to use world's distinctive features

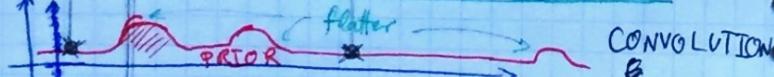
- let's assume our world has 3 distinctive landmarks → 3 doors, all similar and non-distinguishable

- let's assume robot senses a door ⇒ new Belief function



Posterior Belief: the belief after a measurement has been taken!

- now let's assume robot moves forward to the right. Then we can shift the Belief according to motion.



- Now robot measures again a door. Now we multiply our belief, that is prior to the second measurement with the shifted belief.

18

POSTERIOR

- The peak in the resulting function is the only place in the prior that corresponds to the measurement of a door, whereas all the other places of door have a lower prior belief.

- As a result, this function is really interesting.

I + is a distribution that focusses most of its weight on the correct hypothesis of the robot being at the second door.

And it provides very little belief ~~to~~ to places far away from doors and other doors.

Robot localized itself !



Robot
measures
Red

0.04	0.12	0.12	0.04	0.09
------	------	------	------	------

$$\Sigma = 0.36$$

$$) \div 0.36$$

$\frac{1}{9}$	$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{9}$	$\frac{1}{9}$
---------------	---------------	---------------	---------------	---------------

$p(x_i | z)$ posterior distribution

Fig.

- The measurement update function is the absolute key function of localization

 Takes in a measurement and a probability array, and calculates a new probability array based on the measurement.

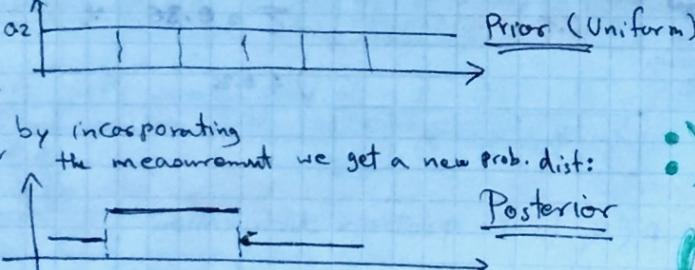
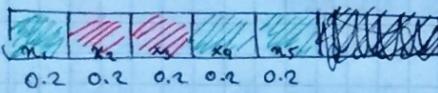
 sense(p, Z):
 $q = []$

$p = [0.2, 0.2, 0.2, 0.2, 0.2]$
 world = [G, R, R, G, G]
 $p_{Hit} = 0.6$
 $p_{Miss} = 0.2$

another way:
 $q = []$
 for i in range(len(p)):
 $hit = (Z == \text{world}[i])$
 $q.append(p[i] * (hit * p_{Hit} + (1 - hit) * p_{Miss}))$
 else
 $q.append(p[i] / p_{Miss})$
 $s = \text{sum}(q)$
 for i in range(len(q)):
 $q[i] = q[i] / s$
 return q

← posterior

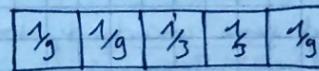
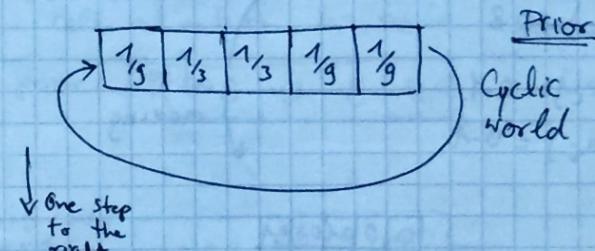
Measurement Update Function
 Update probabilities according to measurement:



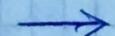
- Now for multiple measurements: we just recursively update our probability vector using new measurements and its prior value:
 $p = [0.2, 0.2, 0.2, 0.2, 0.2]$
 for Z^m in measurements
 $p = \text{sense}(p, m)$

Robot Motion :

- As robot moves, the prior ^{probabilities(?)} shift
 (if the robot movement is perfect: i.e. it moves exactly as far as it thinks it does)

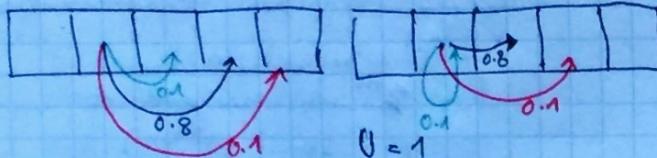


- But, robot movement is inaccurate (not perfect)
 - We have to model this
 - This is the primary reason, why localization is hard.



Inaccurate Robot Motion:

$$U=2$$



- The robot should move $U=2$ cells. It does correctly with $p=0.8$. But with $p=0.1$, it just moves too few or with $p=0.1$ it moves too far, overshooting its target.

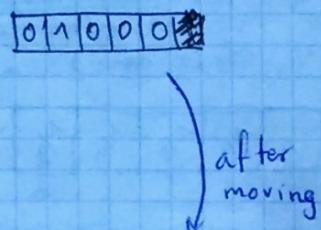
Assume

$$U=2$$

$$P(X_{i+2} | x_i) = 0.8$$

$$P(x_{i+1} | x_i) = 0.1$$

$$P(x_{i+3} | x_i) = 0.1$$



\Rightarrow

- The robot movement added some uncertainty to robot's position!
- Another example

$$\text{prior}$$

0	0.5	0	0.5	0
---	-----	---	-----	---

$$U=2$$

$$\text{posterior}$$

0.4	0.05	0.05	0.4	$\frac{0.05}{0.05}$
-----	------	------	-----	---------------------



Another example: Uniform dist.

$$\text{prior}$$

0.2	0.2	0.2	0.2	0.2
-----	-----	-----	-----	-----

$$U=2$$

$$\text{posterior}$$

0.2	0.2	0.2	0.2	0.2
-----	-----	-----	-----	-----

0.02	0.02	0.16	0.02	0.02
0.02	0.02	0.16	0.02	0.02
0.02	0.02	0.16	0.02	0.02

You can't get any more uncertain than the uniform dist.!

Convolution

Theorem of total probability

Limit (or Stationary) Distribution.

- Assume now that robot moves forever one cell at a time to right.

$$\text{prior: } [1 \ 0 \ 0 \ 0 \ 0]$$

U=1
; forever

$$\text{after many steps: } [0.2 \ 0.2 \ 0.2 \ 0.2 \ 0.2]$$

- As robot continues to get more and more uncertain about where it is, it will eventually reach the state of maximal uncertainty & the uniform distribution

- Every time we move we lose information!

And finally:

The distribution of absolute least information is the uniform distribution.

It has no preference what so ever!

BALANCE

Entropy:

- Entropy represents the amount of uncertainty in a system. Uniform dist is the maximum entropy
- Entropy decreases after the measurement update step and increases after each move() step.

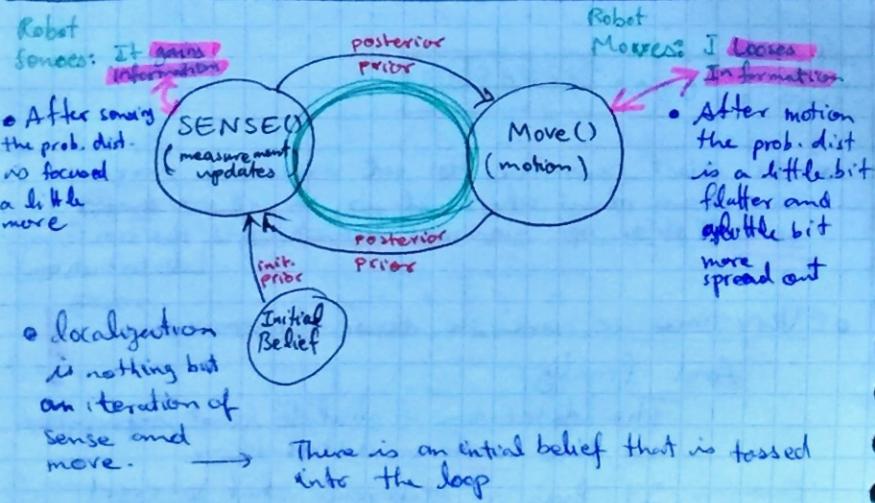
$$\text{Entropy} = \sum (-p \times \log(p))$$

[0.2, 0.2, 0.2, 0.2, 0.2]

⇒ for uniform dis (the max entropy)

$$\text{Entropy} = 5 \times (-0.2 \times \log(0.2)) = 0.699$$

- Taking measurements reduces uncertainty and hence decreases the entropy.
- Let's assume we did a measurement and probabilities became [0.05, 0.05, 0.05, 0.8, 0.05] now the entropy becomes 0.338



- Entropy: Expected log likelihood of the probability of each grid cell

$$\text{Entropy} = -\sum p(x_i) \log p(x_i)$$

- Entropy: a measure of information

Watch the video

Quiz: Sense & move

I + is magic!!!!



Summary:

- Localization maintains a function over all place a robot might be, where each cell has an associated probability value.

BELIEF = PROBABILITY

SENSE = Product of Probabilities times measurements (updating) followed by normalization

Motion = Convolution (=Addition)
(we reversed engineered the question and asked where the robot might have come from, and added these probabilistic

Formal Definition of Localization:

$$\text{Probability} \quad 0 \leq P(x_i) \leq 1 \quad \text{for } x_i \in \text{World}$$

$$\sum_i P(x_i) = 1$$

Measurement: Bayes Rule

Z = measurement x_i = grid cell

$$P(x_i | Z) = \frac{P(Z | x_i) \cdot P(x_i)}{P(Z)}$$

prior
measurement probability normalizer

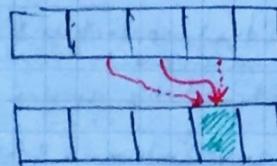
- The measurement update seeks to calculate a belief over my location after seeing the measurement.

$$P(Z) = \sum_i P(Z | x_i) \cdot P(x_i)$$

Motion:

Total Probability

$$P(X_i^t) = \sum_j P(X_j^{t-1}) \cdot P(X_i | X_j)$$



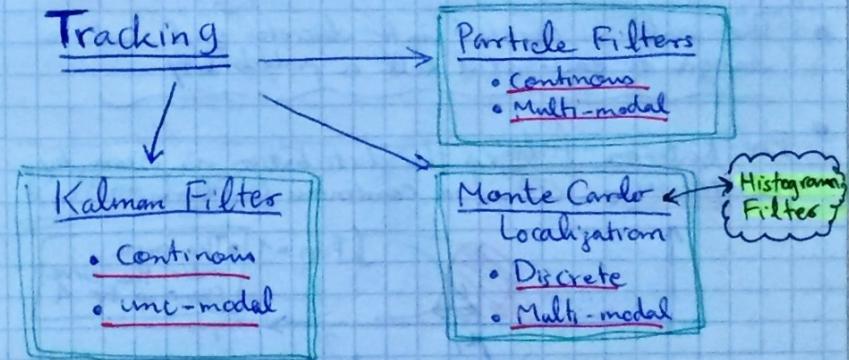
Convolution:
weighted sum over other variable

3. Working with Matrices

Kalman Filter:

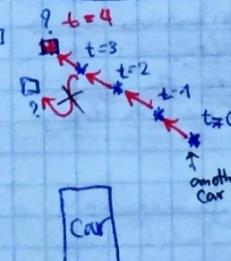
An Algorithm that uses noisy sensors measurements (and Bayes Rule) to produce reliable estimate of unknown quantities (like where the vehicle is likely to be in 3 seconds)

Tracking



- We need to know the position, direction and speed of other cars around us, in order to avoid colliding into them.

- We use the data from laser-range finder for that.



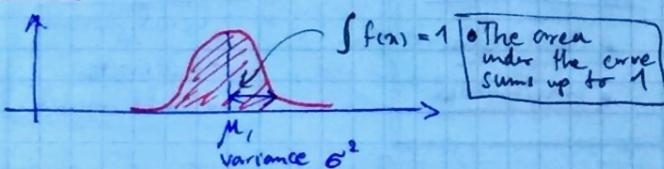
Examples

- We assume a car will go in a straight line, based on the observations $t=0$ to $t=3$.
- Similarly, a Kalman filter takes observations, and estimates future locations and velocities, even when observations are noisy and uncertain.

0.2	0.1	0.5	0.1	0.2
-----	-----	-----	-----	-----



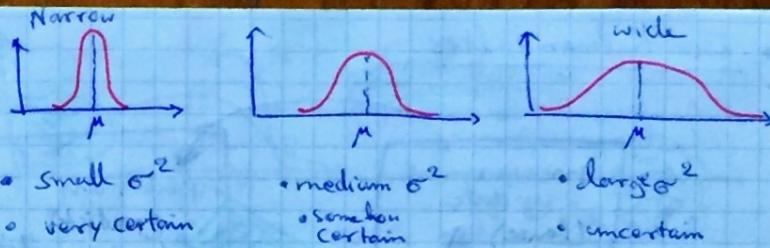
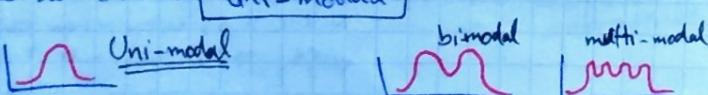
- A histogram is a more approximation of a continuous distribution, by dividing it into discrete partition.
- The world is divided into discrete grid, and we assign to each grid cell a probability.
- In Kalman Filter, the distribution is given by a Gaussian. It is a continuous function.



1-D Gaussian (μ, σ^2)

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp^{-\frac{1}{2} \frac{(x-\mu)^2}{\sigma^2}}$$

- Rather than estimating the entire distribution as a histogram, our task in Kalman filter is to maintain μ and σ^2 that are the best estimate of the location of the object we are trying to find.
- A Gaussian curve has a single peak. It is called "uni-modal".

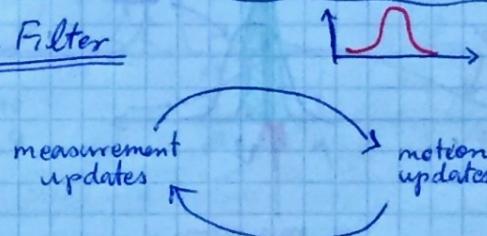


- σ^2 is a measure of certainty.

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp^{-\frac{1}{2} \frac{(x-\mu)^2}{\sigma^2}}$$

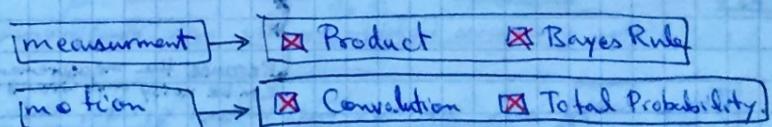
normalizes the difference between x and μ .
so the larger σ^2 , the less important this diff., the less certain the dist., the more wider.

Kalman Filter



- Similar to localization, the Kalman filter iterates measurement updates and motion updates.

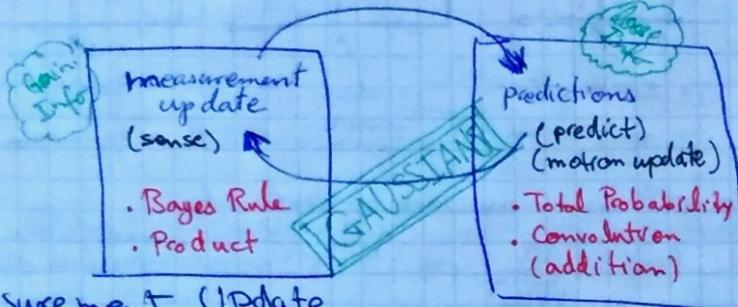
Quiz : In Localization we had



- Measurement meant updating our belief (and renormalizing our distribution).

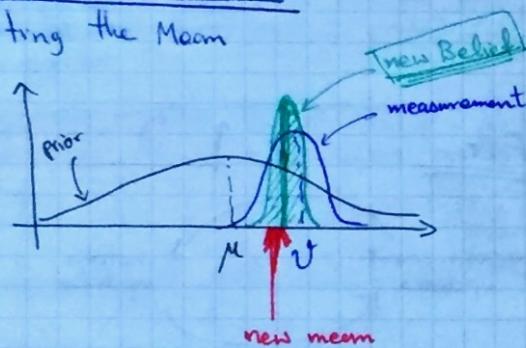
Motion meant keeping track of where all of our probability "went" when we moved (which meant summing up by ^{where we came from} using the law of Total Probability) (29)

In Kalman Filter



Measurement Update

- Shifting the Mean



- The new belief will be somewhere between the previous belief and the new measurement.



- The new belief will be more certain than either the previous belief or the new measurement!

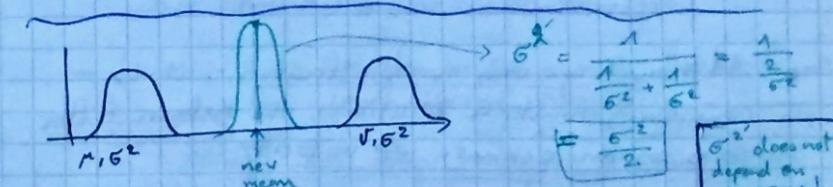
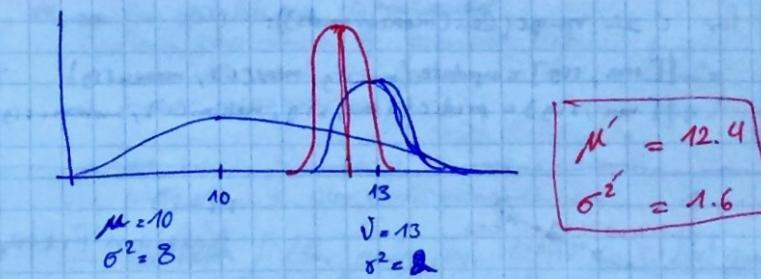
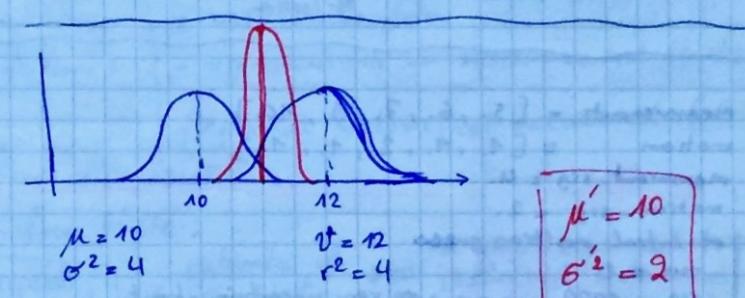
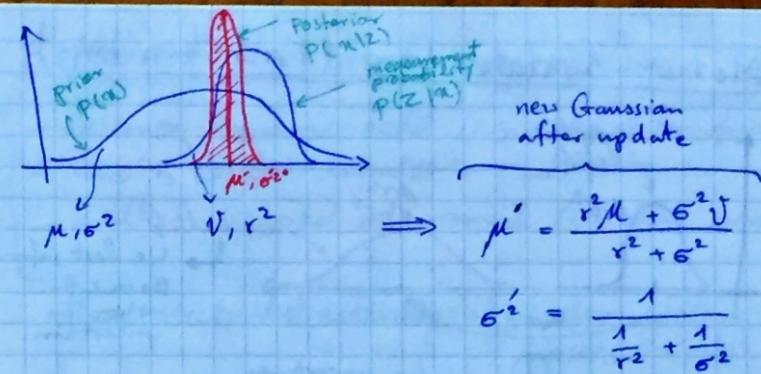
Because the measurement can be noisy and wrong.

Take away lesson: more measurements means greater certainty.

- The resulting Gaussian is more certain than the two component Gaussians.

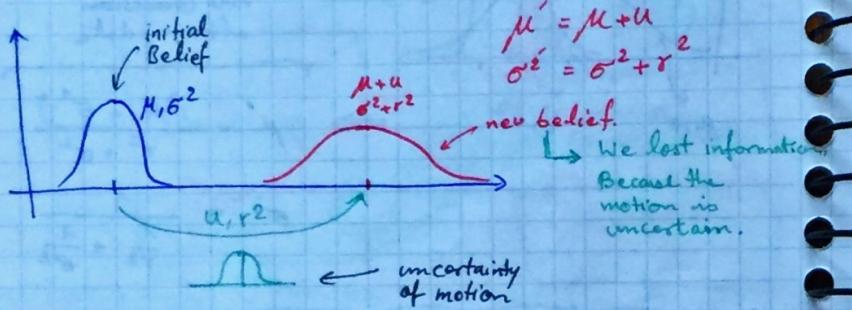
Because we actually gain information!!

The two Gaussians together have higher information content than either one alone.



- Because the Gaussians have the same width, (which means the same certainty), the resulting Gaussian will be in the middle.

Motion Update Step (Prediction Step)



measurements = [5., 6., 7., 9., 10.]

motion = [1., 1., 2., 1., 1.]

measurement-sig = 4.

motion-sig = 2.

initial position guess

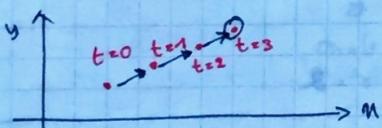
mu = 0

sig = 1000 \rightsquigarrow # very uncertain about our position guess!

1

for i in range(len(measurements)):

(2) [mu, sig] = update(mu, sig, meas[i], meas-sig)
(1) [mu, sig] = predict(mu, sig, motion[i], motion-sig)



n-D
Kalman Filter

- Although we only measure location, but by considering multiple locations, the Kalman filter gives us an estimate of velocity

Kalman Filter [From Wikipedia]

- system's dynamic model (e.g. physical laws of motion)
- known control inputs to that system
- multiple sequential measurements (e.g. from sensors)

\Rightarrow Output: estimate of system state.
(varying variables)

- problems:
- noisy sensor data
 - approximations in system control equations
 - external factors not accounted for

Kalman Filter:

repeat recursively \rightarrow estimate of sys. state = A weighted average of sys's predicted state and the new measurement.

- the weights are determined from covariance of previous state estimate and the covariance of sensor measurements.

- At each time step, we use previous estimate + current measurement
- Kalman Filter only requires the last best guess rather than entire history of sys. state
- Kalman gain: relative weight given to the measurement and current estimate.

gain $\uparrow \Rightarrow$ jumpy, responsive to measurement

gain $\downarrow \Rightarrow$ follow model predictions more closely, smooth out noise, but less responsive to new measurement.

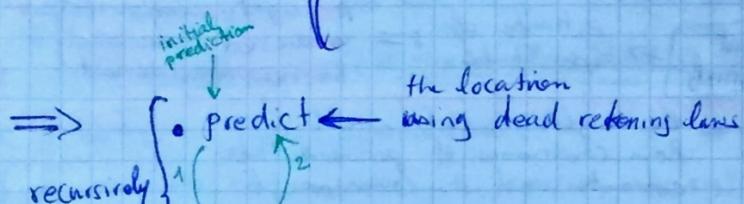
Example 8 Tracking position of a truck

measurement : {

- GPS \rightarrow noisy, jumps within some meters
- GPS covariance: describes the certainty of measurement telling the location

movement model
(Physical model of system) :

- dead reckoning
 - determine new position by direction and wheel rotations and ~~if of~~ ~~velocity~~ (accel?)
 - covariances uncertain at high speed
 - drifts from real position, because of error accumulation



- Use both covariances to weight the combination.

- The GPS ~~at~~ measurement should pull the position estimate back towards the real position, but not disturb it to the point of becoming rapidly jumping and noisy.

For simplicity, assume all matrices are constant,
in reality they also depend on time: $F \rightsquigarrow F_k$

F : the state transition model

H : observation model

Q : covariance of process noise

R : covariance of observation noise

B : (sometimes) control input model

Kalman Filter assumes true state at time k is evolved from previous state at $k-1$ according to:

$$x_k = F x_{k-1} + B u_k + w_k \quad \text{where } w_k \sim N(0, Q)$$

process noise

At time k observation, or measurement of the true state x_k is given by:

$$z_k = H x_k + v_k \quad \text{where } v_k \sim N(0, R)$$

observation noise

Predict

$$\hat{x}_{k|k-1} = F \hat{x}_{k-1|k-1} + B u_k$$

$$P_{k|k-1} = F P_{k-1|k-1} F^T + Q$$

$\hat{x}_{k|k-1}$ = estimate of x at time k , given observation up to and including $m \leq n$

predicted (a ~~priori~~) state estimate

predicted (a ~~priori~~) error covariance

Update

$$g_k = z_k - H \hat{x}_{k|k-1}$$

$$S_k = R + H P_{k|k-1} H^T$$

$$K = P_{k|k-1} H^T S^{-1}$$

$$\hat{x}_{k|k} = \hat{x}_{k|k-1} + K g_k$$

$$P_{k|k} = (I - K H) P_{k|k-1}$$

innovation (or measurement pre-fit) residual

innovation (or pre-fit residual) covariance

optimal Kalman gain

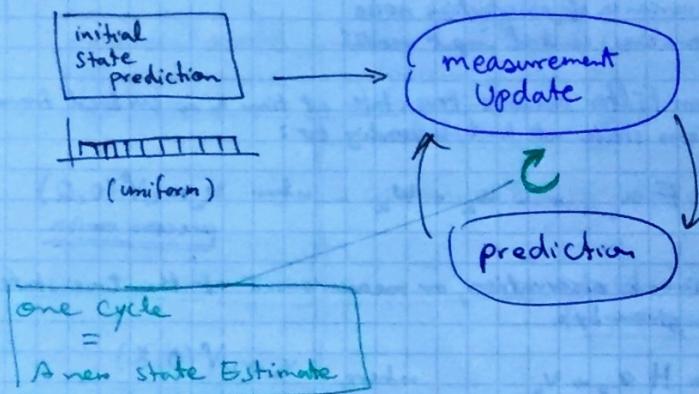
updated (a posteriori) state estimate

updated (a posteriori) estimate covariance

$$\tilde{g}_{k|R} = z_k - H \hat{x}_{k|k}$$

measurement post-fit residual

- The Kalman Filter combines the location sensor measurement with the inaccurate prediction of motion to get a filtered location estimate that is better than any estimate that comes from only sensor or only knowledge about movement.



State $[0, 50 \text{ m/s}] \rightarrow \text{initial state}$

Motion model: $\text{distance} = v \cdot t$

$$\bullet dx = \frac{1}{2} a dt^2 + v_0 dt$$

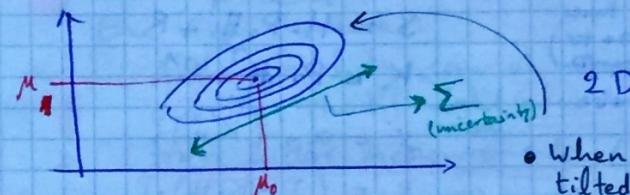
State transformation matrix

Multivariate Gaussians

$$\mu = \begin{pmatrix} \mu_0 \\ \vdots \\ \mu_D \end{pmatrix} \quad \Sigma = \begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix}$$

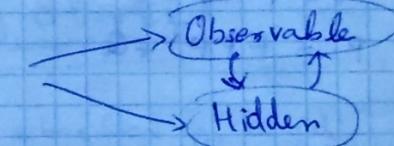
Covariance matrix

$$\text{PDF: } f_X(x_1, \dots, x_D) = (2\pi)^{-D/2} |\Sigma|^{-\frac{1}{2}} \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1} (x - \mu)\right)$$



- When Gaussian is tilted, this means there is correlation between uncertainty of x and y .

Kalman Filter States



Design of a KF:

$$\begin{aligned} \begin{pmatrix} u' \\ i' \end{pmatrix} &\leftarrow \underbrace{\begin{pmatrix} \text{state transformation function} \\ F \end{pmatrix}}_{\text{states}} \begin{pmatrix} u \\ i \end{pmatrix} \\ z &\leftarrow \underbrace{\begin{pmatrix} \text{measurement function} \\ H \end{pmatrix}}_{\text{measurements}} \begin{pmatrix} u \\ i \end{pmatrix} \end{aligned}$$

or state transition function

States

Measurements



PREDICTION

X: estimate

P: uncertainty covariance

F: state transition matrix

u: motion vector

Q: covariance (uncertainty) of process noise

z: measurement

H: measurement func.
(maps states to measurements)

R: measurement noise

$$\hat{x} = Fx + u$$

$$P' = F \cdot P \cdot F^T$$

measurement update

error

Projecting system uncertainty
into the measurement space

$$y = z - H\hat{x}$$

$$S = H \cdot P \cdot H^T + R$$

$$K = P \cdot H^T \cdot S^{-1}$$

Kalman
Gain

$$\hat{x}' = \hat{x} + (Ky)$$

$$P' = (I - K \cdot H) \cdot P$$

Simplified Equations of KF:

$$\hat{x}' = Fx + Bu \quad \xrightarrow{\text{(sometimes) control input model}}$$

$$P' = F \cdot P \cdot F^T + Q$$

$$y = z - H\hat{x}$$

$$S = R + H \cdot P' \cdot H^T$$

$$K = P' \cdot H^T \cdot S^{-1}$$

$$\hat{x} = \hat{x}' + Ky$$

$$P = (I - K \cdot H) \cdot P'$$

$$y = z - H\hat{x}$$

predict
(motion
update)

update
(measure
update)

- Check The Jupyter Notebook on Kalman Filter. This is cool! :)

VERY INTERESTING:

- Kalman Filters can give us insight into variables that we cannot measure directly

Ex. (From notebook):

Although lidars cannot measure velocity directly, the KF can infer the velocity from lidar measurements. (because the motion model (F) models the relationship between distance and velocity)

CPP Getting Started

- typedef `vector<vector<float>>` t-grid;
- function declaration vs. function definition
- #include <iostream> vs. #include "distance.h"
 - double quote looks for the files in the "main.cpp" directory first.
 - Brackets first looks for files where system header files are kept.

L2 - CPP Vectors (2 h 18 min)

- CPP Lists and CPP vectors → Sequence Containers

- using namespace std;
- std::vector<int> myvector (10, 6)
- C++11:


```
std::vector<float> myvector = {5.0, 3.0, 2.7, 8.2, 7.9}
```
- C++98, C++11, C++14, C++17
C++03

L3: Practical CPP: (53 min)

- ISO C++ Standards: C++98, C++03, C++11, C++14, C++17
- g++ -std=c++11 main.cpp
- cin >> integertwo;
- stringstream ifstream matrixfile ("matrix.txt");


```
while (getline (matrixfile, line)) {
    ss >> i; // next value (Also parses!)
    ss.str (line)
    ofstream outfile;
    outfile.open ("matrixoutput.txt");
    outfile << i << " ";
  }
```

L4: C++ OOP (2h 23m)

- include guards
- recommendations
 - ! "put the minimum number #include statements needed into a header file"
 - include guards: prevent redeclaration of things when header files →

← are included

```
#ifndef ENGINE_H
#define ENGINE_H
```

}

```
#endif /* ENGINE_H */
```

• Recommendation:

"Avoid using namespaces in header files.
This ~~can~~ can help avoid naming conflicts"

• std::vector<int>::size_type

- is actually an unsigned int
- is guaranteed to hold up to the maximum size of a vector.

L5: Python vs. C++ Speed

- C++ !
is at least ^①
- Trade-off between
 - speed of execution twice as fast
 - speed of development as Python

- Very fun way of initializing a vector from an array

```
static const arr[5] = {16, 22, 7, 69};
```

```
vector<int> vec (arr, arr + sizeof(arr)/sizeof(arr[0]));
```

pointer pointer

length of array

↓ user the constructor vector(iterator begin, iterator end);

Performance Programming in C++

$16 \times 16 + 2^9 \times 2^9 = 256$

28.8M

L1: Intro to optimization (1h 40m)

- Code optimization :
 - use less memory
 - increase power efficiency
 - execute faster

- ↴ only preprocess, compile, assemble steps

- g++ -C main.cpp
- xxg ~~main.o~~ main.o ← Hex dump
- xxg -b main.o
- std::bitset<8>('a') → 0010 0001
- std::bitset<sizeof(char)*8>(97) → 0010 0001
- float example = 97.0;
- char binary[sizeof(float)];
- memcpy (binary, &example, sizeof(float));
- for (int i=0; i < sizeof(float); i++) {
 std::cout << std::bitset<sizeof(char)*8>(binary[i]);
}

↙ binary representation of 97.0

② Heap vs. Stack:

```
int * pointerVar;  
pointerVar = new int;  
*pointerVar = 10;  
delete pointerVar;  
pointerVar = NULL;
```

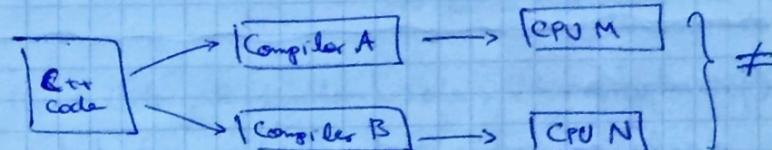
& : address-of operator

* : dereference operator

*foo: (value pointed to by foo)

③ Code optimization:

- Analyzing algorithms
- Understanding how a computer executes instructions
- Learning the nuances of both programming language and compiler you are using.
- ↳ and Hardware



L2: Optimization Practice (2h 30min)

① Hardware:

- ↳ HW Limitations, e.g. CPU's slow software approximation of trigonometric functions
 - ↳ e.g. use small angle approximation if possible.

→ Or, e.g. using 64-bit int on a 16-bit or 32-bit architecture might be inefficient.

② Compiler:

- e.g. compiler unrolls the for loop.
to avoid checking the condition.

[See list of possible compiler optimizations at Wikipedia]

③ Algorithms:

- e.g. quicksort vs. bubble sort
- Check the underlying algorithm when using STL

- fundamental type: (e.g. bool, double, ...)
 - & reference variable \rightarrow a variable always $\&\&$ (since C++11 rvalue reference)
 - rvalue (can't change) (temporary)
lvalue (can change) (variable)
-

Optimization:

- rely on testing and verification instead of instinct!
- Performance testing, Performance analysis (Profiling)
Algorithmic optimization
Platform independent optimization
- SW Platform dependent optimization
HW Platform dependent optimization
- Principle of diminishing returns
- Principle of diminishing portability

• C++ Standard Clock

```
#include <ctime>
Std::clock_t start;
double duration;
start = std::clock();
function_name(var1, var2);
duration = (std::clock() - start) / (double)CLOCKS_PER_SEC;
std::cout << "duration in milliseconds" << 1000 * duration << endl;
```

(44)

• Some Optimization Techniques:

- remove dead code
- avoid extra if statements
- avoid nested for loops
- \rightarrow - avoid creating extra variables
- \rightarrow - reserve space in memory for vectors
- passing variables by reference
- using static \rightarrow use static in functions for values that remain the same between function calls.

$$\left\{ \begin{array}{l} \text{vol} = x + y * z \\ \text{recap} = \sqrt{\text{vol}} \end{array} \right.$$

$$\left\{ \begin{array}{l} \text{recap} = 1 / (\text{x} * \text{y}) \\ \text{vector<int>} \text{ foo}; \end{array} \right.$$

$$\text{foo.reserve(15);}$$

• Dead Code, Redundant Code, Unreachable code

• If statements:

- put most common cases first
- get rid of branches

• gcc -O3 flag \rightarrow 150 ~ 20 ms !

Optimization \rightarrow 350 \rightarrow 150 ms
 \downarrow - O3 : 20 m ! 17 times faster 1700x ◀

Navigating Data Structures

L1: How to Solve Problems: (2 h)

~~Recursive~~

tar -chvfz \rightarrow tar.gz

tar -XVfz \rightarrow extract.

!tar chvfz notebook.tar.gz *

(45)

L3: The search Problem:

Problem Solving:

The theory and technology of building agents that can plan ahead to solve problems.

- We talk about a class of problems where the complexity of the problem comes from the idea that there are many states

e.g. a navigation problem.
↳ sequence of actions



- This is in contrast to e.g. a problem of partial observability, where the actions, paths, results, ... are not fully known.

Formal Definition of Problem:

Initial State $\rightarrow s_0$

Actions (s) $\rightarrow \{a_1, a_2, a_3, \dots\}$
can depend on state

Result (s, a) $\rightarrow s'$

Goal Test (s) \rightarrow True / False

Path Cost ($s_i \xrightarrow{a_i} s_{i+1} \xrightarrow{a_{i+1}} s_{i+2} \dots$)
mostly additive cost $\rightarrow \sum$ step costs

Step Cost (s, a, s') $\rightarrow n$

Cost value (n)
where
 $i = 0, 1, \dots$
 $j = 1, 2, \dots$

- State Space: the set of all possible states
- We navigate the state space by applying actions.
- Frontier: the set of farthest states we have explored.

↳ At any given time we have $\begin{cases} \text{frontier} & \text{Unexplored region} \\ & \text{Explored region} \end{cases}$

- A general function for solving problems:

Tree Search

(superimposes a search tree on the state space)

Function Tree Search (Problem):

frontier = {[initial]}

loop:

if frontier is empty: return FAIL

Path = remove.choice(frontier)

s = path.end

if s is a goal: return Path

expanding} for a in actions(s):

the path

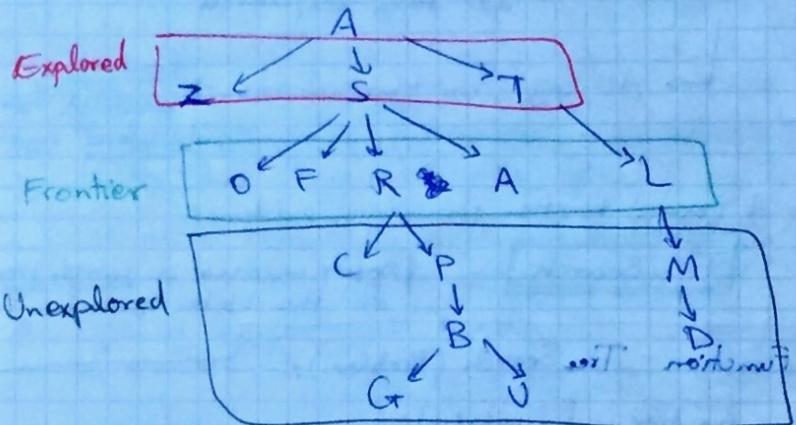
add [path + a \rightarrow Result(s, a)] to frontier

- It is a family of algorithms:

- they all look at frontier,
- pop one off
- look if it is the goal and if yes return it
- Otherwise expand the frontier

- What in tree Search algorithms differs is how to choose a path from frontiers.

- BFS (Breadth-First Search)
(or could be called shortest-first)



• Graph Search

↳ like Tree Search, but don't add states that are in Frontier or are already explored, again to the Frontier.

- BFS is guaranteed to find the path. But shortest path only if the number of step is considered.

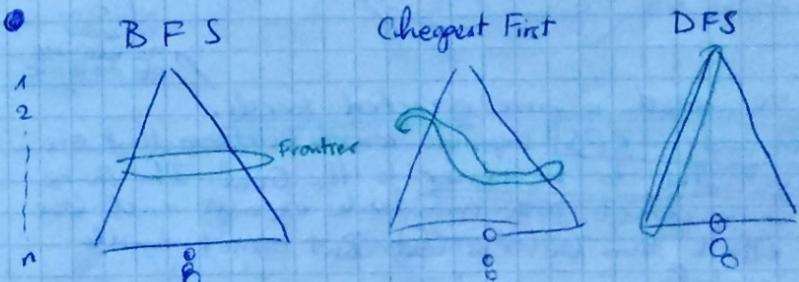
(48)

- Uniform Cost Search ~~Dijkstra~~ (could be called "cheapest first")

- if the cost was not number of steps.

- Cheapest First (Uniform Cost Search) and BFS (or Shortest First) are optimal \rightarrow they are guaranteed to find the shortest path

- But DFS is not optimal, it may find a longer path than the shortest one.



Search Method	Optimal	Frontier size	Size @ n=20
BFS	Yes	2^n	1,048,576
Cheapest First	Yes	2^n	4096, 576
DFS	No	n ↙	20 ↙

→ (49)

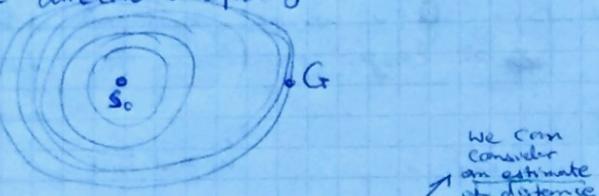
- completeness:

If there is a goal G , will the algorithm find it and terminate?

[Imagine the depth of tree is infinite]

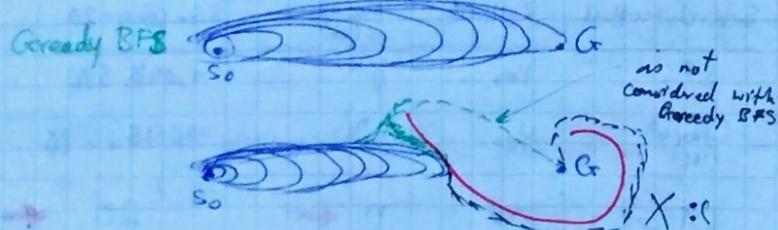
- BFS and Cheapest-First (Uniform Cost) are complete. DFS is not guaranteed to find the solution.

- Uniform Cost Search and BFS expand out in all directions equally



- We need a more directed search.

Greedy BFS expands towards locations that are estimated to be closer to the Goal. If there is a direct path available expands much less effort than Uniform Cost. But cannot handle obstacles.



- We need something that combines both good features:

in many cases explores small number of nodes
→ is directed search towards G and only can handle obstacles and always gives the shortest path

it could be called:

A* Search

A* Best estimated total path cost first

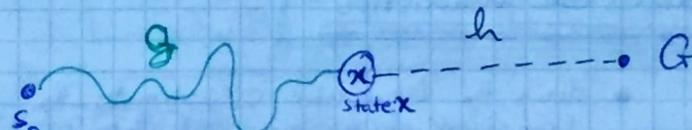
always expand the path that has a minimum value of the function f

$$f = g + h$$

g (path) = path cost

h (path) = $h(s)$ = estimated distance to the Goal

final state of the path



$$f = g + h$$

Cost of the path so far

estimate distance to G

- minimize $g \rightarrow$ keep the path short
- minimize $h \rightarrow$ keep directed towards goal

- A* finds the lowest cost path if:

It depends on the h function

$$h(s) < \text{true cost}$$

- h should never overestimate distance to goal
- h is optimistic
- h is admissible

↳ Wikipedia: Admissible Heuristic

- For an intuition on the proof, see video.
But the idea is, since h is always smaller than actual cost, and since we pick the cheapest first path from frontier, and at the goal $h=0$, then all other paths in the frontier must have a higher cost than the returned path.

i.e.
C =
H =
actual
cost of
P

Admissible heuristics

- does the heuristic overestimate the true cost of reaching goal state?

↳ If yes, it is not admissible

Example

Fifteen Puzzle

1	2	3	4
5	6	7	8
9	11	12	
13	10	14	15

$$h_1 = \# \text{ misplaced blocks}$$

$$h_2 = \text{sum}(\text{distance of blocks})$$

• $h_1 \rightarrow$ Hamming Distance

• $h_2 \rightarrow$ Manhattan Distance

- Both heuristics are admissible!

↳ they are less than actual or equal numbers of moves needed to reach the goal!

↳ Both don't overestimate the actual cost.

- But $h_2 \geq h_1 \Rightarrow h_2$ is better and expands faster (i.e. with fewer steps) and is closer to perfect heuristic

- But how to automatically come up with heuristics?

↳ By relaxing the problem!

- Formal problem/solution rules description

A block can move $A \rightarrow B$
if (A is adjacent to B) relaxation ②
and (B is blank) relaxation ①

• r1 → gives us h_2

• r2 → gives us h_1



- Take the harder problem

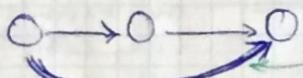
↳ relax ~~one or more of~~ the constraints

↳ you have an easier problem

↳ whose cost is less than or equal the original problem

↳ you have a heuristic that is admissible

- This is like adding new links to the state space.



that can traverse the state space easier.

- adding new operators (links) always makes the problem easier, and thus does not overestimate the cost, and so are admissible.

- When does search not work?
What search can't do?

- Domain is Fully Observable

Search only works if:



- Domain is Known

- Domain is Discrete

- Domain is Deterministic

- Domain is Static

Relaxing the Problem

- Fully Observable: we must see what initial state we start with

- Known

- we must know the set of available actions at each state

- Discrete

- finite number of actions to choose from

- Deterministic

- we must know the result of taking each action

- Static

- nothing else can change the world except our actions.

- Note on Impl:

$$A \rightarrow S \rightarrow F$$

Path is a linked list of Nodes.

Node: { State at the end of path F
Action (S, F)
Cost 234
parent S }



- Frontiers:

Operations: remove best item
add new one
membership test

Impl: Priority Queue

Representative: Set

Build: Hash table
Tree

- Explored

add new member
check for membership

Set

Hash table
Tree

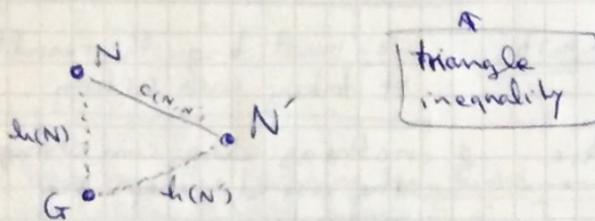
Consistent heuristic

N : node

N' : successor

Monotony

$$h(N) \leq c(N, N') + h(N')$$



- Every consistent heuristic is also admissible. The other way is not true though.

Projects:

- ✓ 1. Joy Ride
- ✓ 2. Histogram Filter
- ✓ 3. Matrix class
- ✓ 4. Translate Python to CPP
- ✓ 5. Optimize Histogram Filter (CPP)
- ✓ 6. Path Planner
- ✓ 7. Line In (July 18th)
- ✓ 8. GritHub (July 21st)
- g. Traffic Light Classifier (August 18th)

From "knowledge" about relations of space
to "know how" to interact properly with space.

- Behavior-Based Robotics \approx Strong Spatial Cognition
- where does these heuristics come from?

Vehicle Motion and Control

L1 Odometers, Speedometers, and Derivatives (1:30 hours)

- Shai Magzimov CEO of Phantom Auto
Ben Shunkman, Software Engineer.

- Odometry: Using data from motion sensors and integrating it in order to come to a coherent belief about how far we have traveled.

- Understanding Motion: position, velocity, acceleration and their relationship.

- Derivatives } Integrals } from Calculus

- Take car's odometry data (e.g. distance traveled) and use it to infer knowledge about velocity & acceleration

- IMU: Inertial Motion unit
 - acceleration along x,y,z
 - rotation rates around x,y,z
(pitch, roll, yaw)

- Vehicle heading and displacement

- Navigation Sensors

- Odometers

Measures how far the vehicle has traveled by counting the wheel rotations.

Useful for measuring distance traveled (displacement)

but are biased because of changing tire diameters

"trip odometer" can be reset by vehicle operator.

- IMU

measures heading, rotation rate, and linear acceleration using magnetometer, rate gyros, accelerometer

• Average Speed : $V_{avg} = \frac{\Delta d}{\Delta t}$

(1:30 h)

L2: Accelerators, rate gyros and Integrals

- Use integral of acceleration ~~area~~ to find velocities and integral of velocities to find displacement.

- Rate Gyros measure angular velocity

Integrating the angular velocity gives us heading

- yaw rate : $\dot{\theta}$

- Real Data is always flawed!

- accelerometers are often biased

↳ careful calibration
↳ better sensors

i.e.
they measure some non-zero value when the acceleration is zero.

- The effect of Error Accumulation

is drastic when integrating over long time intervals ~~area~~ and/or double integrating!

- In order to avoid error accumulation, we should use acceleration data only over short time intervals.