

- IMU: Inertial Motion unit
 - acceleration along x,y,z
 - rotation rates around x,y,z
(pitch, roll, yaw)

- Vehicle heading and displacement

- Navigation Sensors

- Odometers

Measures how far the vehicle has traveled by counting the wheel rotations.

Useful for measuring distance traveled (displacement) but are biased because of changing tire diameters
"trip odometer" can be reset by vehicle operator.

- IMU

measures heading, rotation rate, and linear acceleration using magnetometer, rate gyros, accelerometer

• Average Speed : $V_{avg} = \frac{\Delta a}{\Delta t}$

(1:30 h)

L2: Accelerators, rate gyros and Integrals

- Use integral of acceleration ~~to~~ to find velocities and integral of velocities to find displacement.

- Rate Gyros measure angular velocity

Integrating the angular velocity gives us heading

- yaw rate : $\dot{\theta}$

- Real Data is always flawed!

- accelerometers are often biased

↳ careful calibration
↳ better sensors

↳ they measure some non-zero value when the acceleration is zero.

- The effect of Error Accumulation

is drastic when integrating over long time intervals ~~and~~ and/or double integrating!

- In order to avoid error accumulation, we should use acceleration data only over short time intervals.

Computer Vision and Machine Learning

L1: Computer Vision and Classification (3:41h)

Spatial Coherency:

- As long as the data has spatial coherency we can apply CV techniques to process it
- The data for CV algorithms does not need to necessarily come from cameras.

Spatially Coherent Data:

Any data that predictably varies over space.

Like sound for example. The volume of the sound can give us spatial information. The louder the sound, the nearer the source.

SDCs use in addition to cameras

Sensors, Radars, and Lidars

Radio D. and R. Light Detection and Ranging

• Radars, Lidars → long-range detection

Cameras → rich sensory input

• Radar, Lidars

→ **Active sensors:**

they sense the environment by transmitting energy.

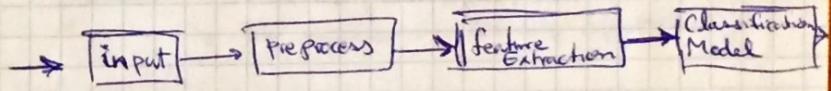
Cameras

→ **Passive sensors:**

can only sense the environment based on energy (photons) already in existing in the scene.

• Part of a good design in a CV system is knowing what not to do, even if it is possible.

Image Classification pipeline



• cropping image img [a:a, b:b, :]

a = height pixels to crop
 b = width pixels to crop

• mask: lower-green = [0, 100, 0]
 upper-green = [100, 255, 100]

mask = cv2.inRange(img, lower-green, upper-green)

masked-img = np.copy(img)

masked-img [mask != 0] = [0, 0, 0]

black
 [0, 0, 0]

61

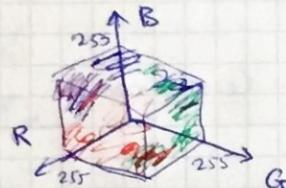
- Masking using simple RGB threshold would not work well under

- variant light conditions
- needs consistent green color

• Color Spaces

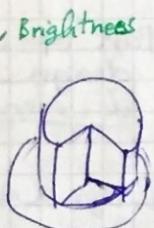
Consider brightness as a separate component

- RGB color space



- HSV color space

- Hue, Saturation, Value



- HLS

- Hue, lightness, saturation

- Because HSV and HLS separate color and brightness into separate channels they are very useful in image processing.

• HSV conversion:

```
lHSV = cv2.cvtColor(img, cv2.COLOR_RGB2HSV)
```

```
h = lHSV[:, :, 0] ← remains consistent under diff. lighting conditions.
s = lHSV[:, :, 1] } vary a lot under diff. lighting condition !
v = lHSV[:, :, 2] }
```

• Classification Labels:

- categorical (e.g. dog, day,..) vs. numerical values
 - integer encoding
 - one-hot encoding

• Two main Types of Features

- in a image:
 - Color-based features
 - Shape-based features

• High-pass filter:

Detect big changes in intensity over a small area.

0	-1	0
-1	4	-1
0	-1	0

↑
High-pass filter
Kernel

- The weights sum to zero
- The center weight is the highest
- The neighbor weights are negative

Detect Edges

Edge: Where intensity changes abruptly

- The edge detection filter computes the difference (or change) between neighboring pixels.

It is an approximation of Derivative of
Image over space

- A kernel for detecting horizontal lines

$$\rightarrow \begin{pmatrix} -1 & 0 & 1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{pmatrix}$$

- finds the difference between the top and bottom edges surrounding a given pixel.

- Sobel Filter: for edge detection

$$S_x = \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix}$$

Taking an approx. of

$$S_y = \begin{pmatrix} -1 & -2 & 1 \\ 0 & 0 & 0 \\ +1 & 2 & 1 \end{pmatrix}$$

Image Derivative

- Frequency in an Image is the Rate of change

- Images change in space. A high frequency image is where the intensity changes a lot, and the level of brightness changes quickly from one pixel to another.
- A low frequency image is a one where brightness changes slowly or is relatively uniform.
- Most images have both high frequency and low frequency components.
- High frequency components correspond to edges of the object in an image

Roberts Cross Operators:

(AKA: Roberts Cross)

- Simple, quick to compute, 2-D spatial gradient measurements on an image
- Highlights the regions of high spatial frequency, which often correspond to edges.

$$G_x = \begin{pmatrix} +1 & 0 \\ 0 & -1 \end{pmatrix} \quad G_y = \begin{pmatrix} 0 & +1 \\ -1 & 0 \end{pmatrix}$$

- diagonal gradients. Designed to respond maximally to edges at 45°
- due to small kernel size, susceptible to noise

Histograms :

hsv = cv2.cvtColor(rgb-img, cv2.COLOR_RGB2HSV)

h-hist = np.histogram(hsv[:, :, 0], bins=32, range=(0, 180))

s-hist = np.histogram(hsv[:, :, 1], bins=32, range=(0, 256))

v-hist = np.histogram(hsv[:, :, 2], bins=32, range=(0, 256))

bin-edges = h-hist[1]

bin-centers = (bin-edges[1:] + bin-edges[0: len(bin-edges)-1]) / 2

fig = plt.figure(figsize=(12, 3))

plt.subplot(131)

plt.bar(bin-centers, h-hist[0])

plt.xlim(0, 180)

plt.title("H Histogram")

rest two graphs for S and V histograms

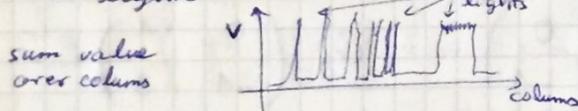
- Spatial Domain
- Frequency Domain

• Summation to create Feature Vectors

- To keep spatial information it is also useful to sum up pixel values alongs columns or rows.

⇒ Spike in various colors or intensity values over space.

Example: A night image has many ^{often} high intensity lights.



Self Driving Car Engineer :)

Projects:

- ✓ - Finding Lane Lines
- ✓ - Advanced Lane Finding
- ✓ - Traffic Sign Classifier
- ✓ - Behavioral Cloning
- ✓ - Extended Kalman Filters
- Kidnapped Vehicle
- Highway Driving
- PID Controller
- - Linked In Profile
- - GitHub Profile
- Programming a real Self Driving Car

Part 1: Computer vision, Deep learning, Sensors (EKF, UKF) Classification Behavior cloning (End-to-End Learning)

Part 2:

Localization, Path Planning, Control, Sys. Integration

Lesson 1 : Introduction

Sebastian Thrun

- 2003 DARPA Challenge (Stanley)
- 2009 the Stanley team joined Google

- David Silver (worked at Ford)
- Ryan Keenan (astro physicist)

- Two distinct approaches to Autonomous Development:

- Robotic Approach

Fuse sensors to measure vehicle's surroundings and then navigate accordingly

- Deep learning

learn to mimic human behavior using deep learning.

1. Content

Computer Vision, Deep Learning,
~~Sensor Fusion~~

2. Localization, Path Planning,
Control, System Integration

Projects :

P1 • Find lane Markings

Lane Markings tell you where to drive

P2 • Advanced lane finding

P3 • Traffic Sign Classifier

P4 • Behavioral Cloning

clone human behavior using deep learning on camera images

P5 • Extended Kalman Filter

Sensor Fusion

P6 • Kidnapped Vehicle localization

P7 • Path Finding

path planning : come up with a sequence of actions. City navigation, parking lot navigation

P8 • PID controller

Control : steer the actuators to meet an objective

P9 • Put your code in a real SLD car

①

②

63

Extra Circular

• Additional Content

Unsupervised Kalman Filters, Model Predictive Control, Advanced Deep Learning / Semantic Segmentation, Functional Safety

• Autonomous Systems Interview

• Technical Interview Prep

Career Portal

- Career track (tips)
- Career Services

- Feedback on Resume, Cover letter, LinkedIn, Githubs
- Career track tips, informational interviews
- Udacity profile for Talent Program.

Udacity Talent Program

- graduate
- complete Udacity Profile
- Upload Resume and make it public.

A Timetable for Studying and Objectives in the next days [I don't download the course material for the time being for some time]

• L4: Computer Vision Fundamentals (3 hr 46 m)

• P1: Finding lanes 28-30. March 2020 (2 days)

• L6: Camera Calibration (2 hours) 31. March

• L7: Gradients and Color Spaces (2 h) 01. April

• L8: Advanced CV (2 h) 02. April

• P2: Advanced Lane Finding 03. April - 06. April (3 days)

• L10: Neural Networks (3 hr) 07. April 22:00
09. April 22:00

• L11: Tensor Flow (2 h) 09. April 22:00
11. April 21:30

• L12: Deep NN (2 h) 11. April 21:30 - 12. April 20:15

• L13: CNNs (2 h) 12. April 20:15 - 14. April 23:30

• L14: LeNet for traffic signs (2 h)
15. April

12th • P3: Traffic sign classifier 15. April 21:30
24. April 23:30
26. April (9 days)

• L16: Keras (2 h) 26. April 18:30 - 20:30

• L17: Transfer Learning (1 hr 30)
26. April 20:30 - 27. April 23:59

15th • P4: Behavioral ~~Modeling~~ Cloning
28. April
05. May
(7 days)

• L20: Sensors (0:30 m) 06. May 23:00-00:30

• L21: Kalman Filters (0:32 m) 06. May 21:00
08. May 21:00

• L23: Geometry & Trigonometry Review (8 min)

L24: Extended Kalman Filters (5 days) May 09 22:00
May 14 21:00

19th P5: Extended Kalman Filters Project (4 hours) May 14 21:00
May 18 20:00

[Localization, Path Planning, Control, and System Integration]

L1: Intro to Localization (19 min) 19. May 18:00

L2: Markov localization (3 h) 20. May 12:00
22. May 20:00

L3: Motion Models (30 min) 21. May

L4: Particle Filters (3 hours) 24. May 21:30
25. May 20:30

L5: Implementation of a particle filter (1 h) 26. May

P6: Kidnapped Vehicle 27. May - 28. May 11:00

26th

L7: Search (1 h 30 m) May 28./29. May

L8: Prediction (1 h 30 m) 01. June
02. June 17:00

L9: Behavior Planning (2 h) 02. June 18:00
03. June 20:00

L10: Trajectory Generation (2 h) 03. June
(5 days) 08. June

03 May P7: Highway Driving 09. June
22. June (14 days)

L12: PID control (2 h) 26. June

P8: PID controller 29. June 19:30
30. June

{ GitHub, Linked In }



← L16: Autonomous Vehicle Architecture (30 m)
01. July

L17: Introduction to ROS (1 h) 01 July

L18: Packages and Catkin Workspace (1 h)
02. July 2020
22:00

L19: Writing ROS nodes (1 h) 03. July

P9: Program on Autonomous Vehicle 08. July
14. July

15th May L21: Complete the Program 15. July

L4: Computer Vision Fundamentals (3h:46m)

- Feature that could help in identifying lane lines in an image of a highway
 - Color, shape, orientation, Position in Image
- image thresholding
- Region ~~eff~~ Masking

masking the region of interest in the image under the assumption that the camera is placed at a fixed position in ~~teleca~~ front of the car and

⇒ the lane lines will always appear in the same general region of the image

- np.polyfit() $\Rightarrow A_n \cdot x + B$

$$\begin{array}{c} A_1x \\ A_2x \\ A_3x \end{array} \rightarrow \begin{array}{c} A_1x + B_1 \\ A_2x + B_2 \\ A_3x + B_3 \end{array}$$

- np.meshgrid() $\Rightarrow XX, YY$

region-thresholds = $(YY > (A_1 \cdot XX + B_1))$

$\& (YY > (A_2 \cdot XX + B_2))$

$\& (YY < (A_3 \cdot XX + B_3))$

• Both masks

result-image [~color-thresholds & region-thresholds]
 $= (255, 0, 0)$

draws the lane lines in the ROI red.

• Canny Edge detection

- gray \rightarrow gradient \rightarrow threshold for following the gradients.

$$\begin{aligned} \text{image: } f(x,y) &= \text{pixel value} \\ \frac{df}{dx} &= \Delta(\text{pixel value}) \end{aligned}$$

- find edges by tracing out the pixels that follow the strongest gradients.

- rapid changes in brightness in a gray scale image are edges.

- gradient $\frac{df}{dy} = \Delta(\text{pixel value})$

(a measure of change in the image function)

- computing the gradient (derivative with respect to x and y) gives us thick edges.
- Canny Edge detections gives us thin edges as pixels
- Steps:

- convert image to gray
- Gaussian Blur : suppress noise and spurious gradients by averaging
- kernel_size = 3
 $\text{blur_gray} = \text{cv2.GaussianBlur}(\text{gray}, (\text{kernel_size}, \text{kernel_size}))$

- run Canny

$\text{cv2.Canny}(\text{blur_gray}, \text{low_thresh}, \text{high_thresh})$

- ratio between $\frac{\text{low_threshold}}{\text{high_threshold}} \approx 1:2$ or $1:3$

- - gradient intensities over high-thresh are "sure edges"
- gradient intensities below low-thresh are "none edges"
- for each point in between, do connectivity check: keep those that are connected to edges

• Paul Hough $y = mx + b$

- image space line $y = m_0x + b_0$

- Hough space (m_0, b_0)

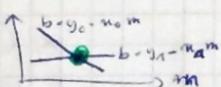
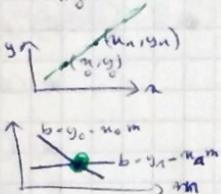
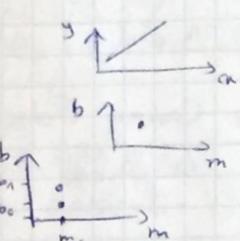
- parallel lines in Hough space

- A point in ~~Hough~~ (x_0, y_0)

image space corresponds to a line in Hough space

$$b = y_0 - x_0 m$$

$$b = -x_0 m + y_0$$



~~Hough~~

- An intersection point of two lines in Hough space corresponds to a line passing through two points in image space.

- Strategy to find line in image space, find intersection of lines in Hough space.

\Rightarrow • First do Canny to detect edges.

- Every point on an edge corresponds to a line in Hough space

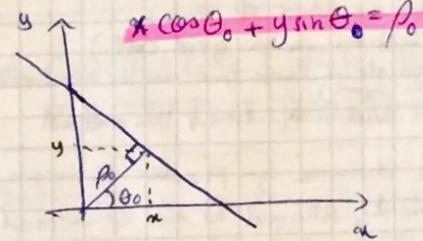
- where many lines intersect in Hough space, we have found a line in image space that connects all those points.

Problem! Vertical lines have infinite slope.
in (m, b) representation.

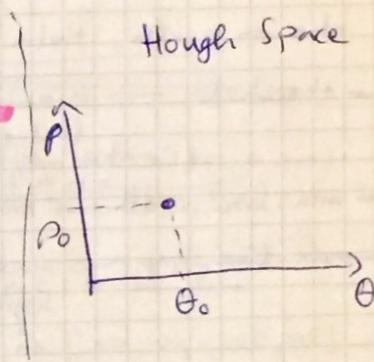
$$y = m n + b$$

\Rightarrow Let's represent a line in image space in polar coordinates

Image Space



Hough Space



- this representation has the advantage that we can represent any line (vertical lines too).

- now each point in image space corresponds to a sin curve in Hough space

- A line of points, translates to a whole bunch of sin curves in Hough space

\Rightarrow find intersection of ~~sin~~ curves

sin

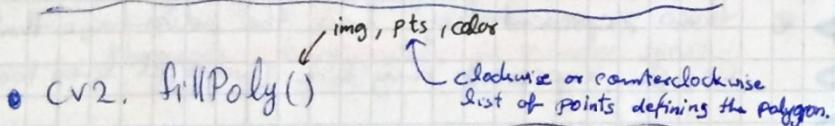
curves

- cv2.HoughLinesP (masked_edges, rho, theta, threshold,

- rho } distance
- theta } and
angular
resolution of
the grid $\frac{1}{10}$ in Hough space
Search for intersections

rho in pixels, theta in radians

- threshold \rightarrow minimum number of votes
(intersection in a grid cell)
(i.e. how many points in image space must be associated with the line segment)
- min_line_width \rightarrow min width of lines in output
- max_line_gap \rightarrow max distance in pixels between segments allowed to be connected in a single line.

cv2.fillPoly()


• Parameter Tuning \rightarrow a big challenge!

28-30.07.20

P1: Finding Lane Lines

- Gray, blur
- Canny
- ROI Mask
- Hough

- Group segments
- fit lane line based on segments
- draw lane lines

- cv2.fitLine (points, DIST, param, r-eps, a-eps)

cv2.fitLine (mid-points, cv2.DIST_L2, 0, 0.01, 0.01)

returns a vector colinear to the line
a point on the line : $[vx, vy, x0, y0]$

the line equation

$$(x, y) = (x_0, y_0) + t(v_x, v_y) \text{ for } t \text{ from } -\infty \text{ to } \infty$$

we need to find t in image boundaries
to clip the line :

$$t_0 = (0 - y_0) / v_y$$

$$t_1 = (\text{line-height} - y_0) / v_y$$

$$p0 = (\text{line}[2:4] + t_0 * \text{line}[0:2]).\text{astype(np.uint32)}$$

$$p1 = (\text{line}[2:4] + t_1 * \text{line}[0:2]).\text{astype(np.uint32)}$$

cv2.line (img, P0, P1, thickness, color)

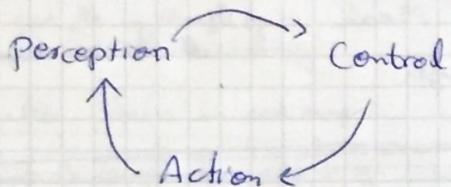
29

draw the line

L6: Camera Calibration

31.03.20

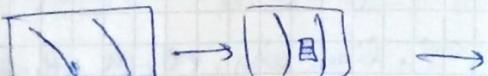
Robotic Cycle:



- 80% of the challenge of building a SDC is **perception** — Sebastian Thrun.
- Why use cameras over Radars and Lidars?

Sensors	Spatial Resolution	Cost	3D
Radar + Lidars	Low	\$\$\$\$	Yes
Single Camera	High	\$	No

- To steer the car we need to know how much the **lane** is curving.
→ transform perspective to above

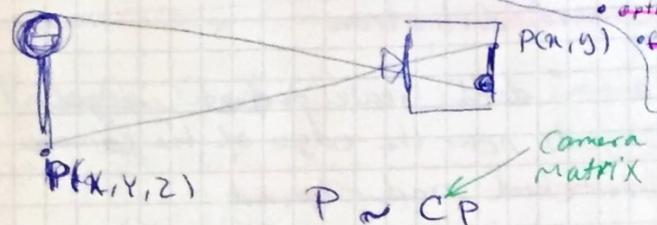


- But to do that we have to first fix the **camera distortion**
- Cameras don't create perfect images. Objects near the **edges of the camera** are **stretched** and **skewed**
- IMPORTANT:**
If the lens is distorted, we will get the wrong measurement for **lane curvature**. So our steering commands will be wrong!
- Distortion results in
 - change of apparent geometry of objects
↳ size, shape
 - make objects appear closer or farther away
 - make object appearance change depending on where the object is in the field of view.

Pinhole Camera Model

X :)

Pinhole Camera Model



$$3D \rightsquigarrow 2D$$

Camera
matrix

- image plane Π
- principal point c
- center of projection
- optical axis

• focal length

- Most distortions can be ~~eliminated~~ captured by 5 numbers: the distortion coefficients

$$(k_1, k_2, p_1, p_2, k_3)$$

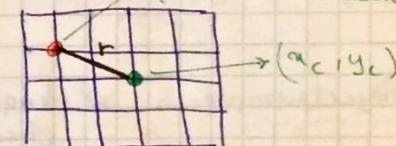
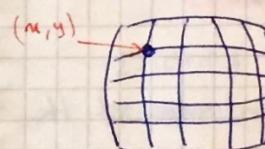
tangential radial

represent radial and tangential ~~distortions~~ distortions in an image.

- sometimes even more coefficients are required to describe the distortion, in severe cases.

- distortion \rightarrow Cam. Calibr. \rightarrow undistorted Image.

- Distortion correction: $\rightarrow (x_{\text{corr}}, y_{\text{corr}}) \rightarrow (x_c, y_c)$



(x, y) a point in distorted image

$(x_{\text{corr}}, y_{\text{corr}})$ same point in corrected image

r : distance of the corrected point to center of image distortion, which is often center of the image. (x_c, y_c) is also called distortion center.

k_3 : is required to correct major distortions and is often near zero for normal cameras. We take k_3 also into account.

Radial Distortion correction:

$$x_{\text{distorted}} = x_{\text{corr}}(1 + k_1 r^2 + k_2 r^4 + k_3 r^6)$$

$$y_{\text{distorted}} = y_{\text{corr}}(1 + k_1 r^2 + k_2 r^4 + k_3 r^6)$$

Tangential Distortion Correction

$$x_{\text{corr}} = x + [2P_1xy + P_2(r^2 + 2x^2)]$$

$$y_{\text{corr}} = y + [2P_2xy + P_1(r^2 + 2y^2)]$$

- We need a **known**, regular, high contrast pattern for calibration (correcting for the distortion effect)

\Rightarrow Chessboard

(img, patternSize, flags=None)

- cv2.findChessboardCorners()
- cv2.drawChessboardCorners(img, patternSize, corners, ret)
- It is recommended to use **at least 20 images** to get reliable calibration.

Object-points

[0,0,0]

known coordinates of the chessboard
3D
(8 rows)
[0, 1, 2, 3, 4, 5, 6, 7]

to [0, 1, 2, 3, 4, 5, 6, 7] for a 7x5 chessboard

the z coordinate of all

object points will be zero, because they all lie on the same flat surface.

$$\text{objp} = \text{np.zeros}((7*5, 3), \text{np.float32}) \begin{bmatrix} 0 & 0 & 0 \\ 1 & 0 & 0 \\ 2 & 0 & 0 \\ \vdots & \vdots & \vdots \\ 6 & 0 & 0 \end{bmatrix} \quad 35 \times 3$$

$$\text{objp}[:, :, 2] = \text{np.mgrid}[0:7, 0:5].T \cdot \text{reshape}(-1)$$

$$\text{impoints} = \# \text{corners} \quad \longleftrightarrow \quad 35 \times 2$$

- glob \rightarrow read images with a given filename pattern

- cv2.calibrateCamera (obj points, im points, size of the image, gray.shape[:,-1], none, none)

returns ret, mtx, dist, rvecs, tvecs
(intrinsic) camera matrix
distortion coeffs
rotation vectors
translation vectors

- cv2.undistort (img, mtx, dist.coeffs, None, rvecs, tvecs)

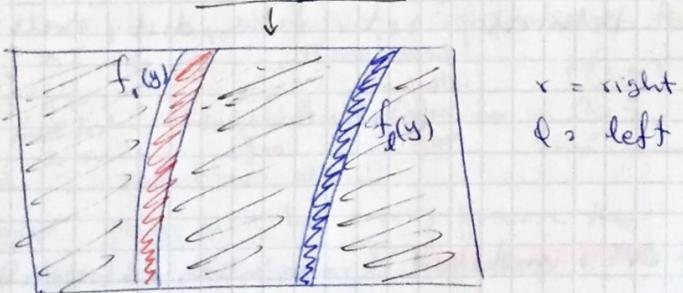
\rightarrow destination image
(undistorted)

Lane Curvature

- We can calculate the **steering angle** with information about speed and dynamics of the car and **how much the lane is curving**.

Steps :

- detet the lane lines using masking and thresholding
- perform a **perspective transform** to get a **birds eye view** of the lane **Rectify the image**
- Fit a (2D ?) polynomial to the lanes in **Birds eye view**.



$$f_r(y) = A_r y^2 + B_r y + C_r$$

$$f_l(y) = A_l y^2 + B_l y + C_l$$

Perspective :

- We can characterize perspective using 3D coordinates with objects with **larger Z** (^{depth} distance) with respect to camera, appear smaller in the 2D image
- A **Perspective Transform** essentially changes the apparent z coordinate of the object points, to change the 2D representation of the object.
- A perspective transform warps the image and effectively drags points towards the camera or pushes them away.
- A perspective transform lets us change our ~~view~~ ^{perspective} to view a scene from different view points and angles:
 - side of the camera
 - below
 - above
- We do perspective transform ~~using~~ using correspondence points from two images
 - linear transformation
 - from one image to the other.

- select 4 points that define a rectangle on the original image (input image, source image)

→ 4 Points are enough to define a linear transform

M =

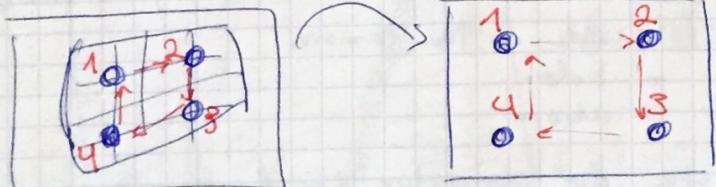
cv2.getPerspectiveTransform(src, dest)

cv2.warpPerspective(img, M, img_size, cv2.INTER_LINEAR)

interpolation method for missing points

- Select 4 points that make a rectangle in SRC image

- Select/define their corresponding 4 points in dest image



- color and Gradient thresholds, a way for determining where certain objects occur in images that have already been undistorted and perspective transformed.

L7: Gradients and Color Spaces

01.04
20

- Canny gives us a lot of edges on scenery and surrounding cars that we end up discarding.

- The lane lines are close to vertical. This ~~is~~ we know a priori.

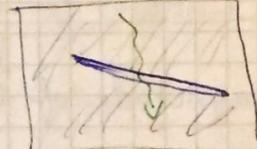
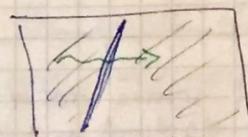
⇒ Take advantage of this fact by using gradients smarter to detect steep edges.

- Canny takes derivative w.r.t. x and y
- Sobel operator

$$S_x = \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix} \quad S_y = \begin{pmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{pmatrix}$$

gradient in x direction

gradient in y direction



- gradient in x direction emphasizes more the vertical lines.
- gradient in y direction emphasizes more the horizontal lines.

(89)

gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)

sobelx = cv2.Sobel(gray, cv2.CV_64F, 1, 0)

sobely = cv2.Sobel(gray, cv2.CV_64F, 0, 1)

Absolute value of Sobel x

abs_sobel = np.abs(sobelx)

Convert the abs value image to 8 bit

scaled_sobel = np.uint8(255 * abs_sobel / np.max(abs_sobel))

Create a binary threshold to select pixels based on gradient strength

thresh_min = 20

thresh_max = 100

sxbinary = np.zeros_like(scaled_sobel)

sxbinary[(scaled_sobel >= thresh_min) &

(scaled_sobel <= thresh_max)] = 1

plt.imshow(sxbinary, cmap='gray')

Magnitude of Gradient : square root of sum of gradient in both directions

$$\text{abs_sobel} = \sqrt{(\text{sobelx})^2}$$

$$\text{abs_sobely} = \sqrt{(\text{sobely})^2}$$

$$\text{abs_sobelxy} = \sqrt{(\text{sobelx})^2 + (\text{sobely})^2}$$

Direction of Gradient

$$\arctan\left(\frac{\text{sobely}}{\text{sobelx}}\right)$$

Each pixel in gradient image has a value for gradient direction between $-\frac{\pi}{2}, \frac{\pi}{2}$

Orientation of 0 \Rightarrow vertical line

Orientation of $\pm \frac{\pi}{2} \Rightarrow$ horizontal line

np.arctan2 \Rightarrow returns values between $-\pi, \pi$

but we take abs value of sobely

$\Rightarrow n \geq 0 \Rightarrow \arctan2 \Rightarrow$ returns between $0, \pi$

Difference atom() and atom2()

only quadrant I and IV



all quadrants

the result of direction of Gradient is very noisy though

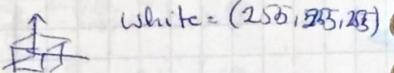


Color Thresholding:

- RGB thresholding does not work well in ~~when~~ variable light conditions or when lane lines have different colors (like white & yellow)

• Color Space

- RGB



- HSV Hue, Saturation, Value



- HLS

L = lightness

$$0 \leq H \leq 179$$

• Advantage of HLS

- isolates the lightness (L) component which varies the most under different lighting conditions
- Hue and Saturation channel stay fairly consistent in shadow or excessive brightness.

• Combining color and gradient thresholding with ~~the~~ S channel.

L8 : Advanced Computer Vision

01.04.20
21:00

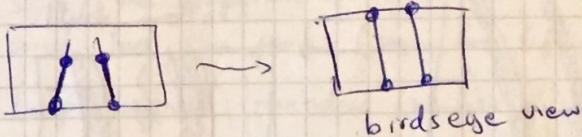
We have covered so far:

- cam. calibration
- distortion correction
- Color and gradient threshold
- ~~Perspective Transform~~

Then we will do

- Detect lane lines
- Determine the lane curvature

- We first do camera calibration (only once)
- Every frame gets distortion corrected using dist-coeffs.
- We then apply thresholding to get the lane lines clearly visible
- We then do perspective transform. We assume that the road is flat. We choose four points lying on 2 straight lane lines.



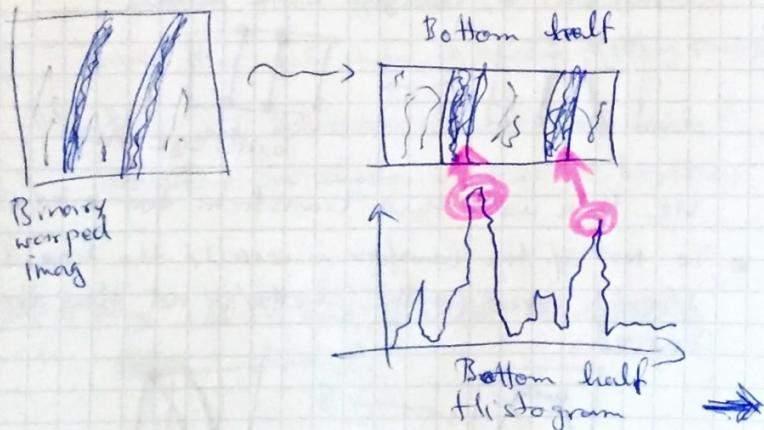
- We then use this transform for every frame
- To test if the transform is correct, the lane lines should appear parallel, whether or not they are curved.



Finding the Lanes: Histogram Peaks

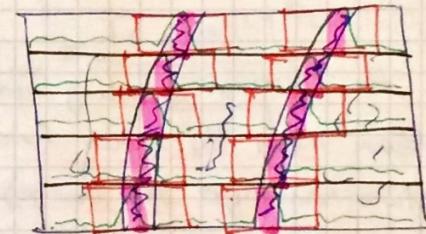
- input: Thresholded warped image
 - there are many ways to do finding the lanes
 - one method is **Peaks in Histograms**
 - We need to determine:
 - which pixels are part of the lines
 - which belong to the left line and
 - which belong to the right line
- ⇒ possible solution: Plotting the histogram of binary activations in image

take only the bottom half,
Lane lines are likely to be
most vertical nearest to the
car.



Sliding Window

- The two peaks in the histogram are good indicators for the x-position of the base of the lane lines
- We use that starting position as the starting point of our search, for lines.
- We use a **sliding window** placed around the line centers to follow the lines up to the top of the frame.



Sliding Window

Histogram in each SW gives us the position of that segment of lane.

Steps :

- split the histogram at base into two (left,right) and find the peaks (bases for left & right)
- Loop through window (number of windows)
 - find the boundaries of window (We don't take the whole width, but a small window; and count all the pixels inside it; If number of pixels is larger than a minpix threshold, we move the window to the mean position of these pixels)
 - Add the (x,y) coordinates of nonzero pixels in the left and right lane lists

Fit the polynomial, np.polyfit()

- After tracking (finding) the lanes in one frame using sliding window, we don't need to do the sliding window in every consecutive frame.

- Instead we create a search margin around the found lanes and only search inside that margin.

- Of course the current sliding window algorithm has a problem when the lanes have sharp curvature and go off the right or left edge of the image. This is related to the "minpix" threshold to move the position of the next sliding window.

Measuring Curvature

- What we have now is a two second order polynomials of lane lines.

$$f(y) = Ay^2 + By + C$$



- We fit for y instead of x , $f(x)$, because $f(y)$

these lane lines are nearly vertical, and if we fit for $f(x)$ chances are that there is same x value for more than one y value.

$$K \in \text{Curvature} \quad R: \text{radius of curvature} = \frac{1}{K}$$

- At point x of the function $f(x)$ the radius of the curvature is given by

$$R_{\text{curve}} = \frac{\left[1 + \left(\frac{dy}{dx}\right)^2\right]^{3/2}}{\left|\frac{d^2y}{dx^2}\right|}$$

$$f(y) = Ax^2 + Bx + C$$

- We have a sec. order polynomial \Rightarrow

$$f'(y) = \frac{dy}{dx} = 2Ax + B$$

$$f''(y) = \frac{d^2y}{dx^2} = 2A$$

$$\Rightarrow R_{\text{curve}} = \frac{\left(1 + (2Ay + B)^2\right)^{3/2}}{|2A|}$$

- For the curvature of the lane at the nearest point to the car, put $y = \text{img.shape[0]}$ in the formula.

- The radius we calculated is in pixel space
- We need to convert it to real-world space.

↳ For this we require to know how long and wide is the section of the road that we are projecting in our warped image

→ we assume the sections for the images in this course that we are projecting is 30 m long and 3.7 meters wide

→ American U.S. road standard

- minimum lane width = 12 ft = 3.7 m
- dashed lane lines = 3 m long
10 ft

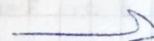
① Let's say camera image

is 720 pixel in y-dimension and we have roughly 700 relevant pixels in x-dimension.

$$\Rightarrow \text{y m-per-pix} = 30 / 720 \quad \# \text{m/pix in y-dim}$$

$$\text{x m-per-pix} = 3.7 / 700 \quad \# \text{m/pix in x-dim}$$

~~P2: Action~~



02.04.20
10 p.m.
03.04.20

P2: Advanced LaneFinding

- More realistic highway scenario

- lighting changes
- color changes
- under passes

- If you do well in this scenario you are reaching state of the art in CV on SDCs

— Sebastian Thrun

- Expected ^{0.0f} magnitude of curve $\approx 1 \text{ km}$

- Assumption: camera is mounted at the center of the car

\Rightarrow mid point at the bottom of the image is between two lines (lane center)

- check offset of the lane center from center of the image (Converted from pixels to meters) is your distance from center of the lane.

- Define a Line() class to keep track of detected line information

- Sanity check: Detected lanes

- have similar curvature
- are separated by approx. right distance
- are roughly parallel.

- Look Ahead Filter: search from prior in consequent frames.

use ↑

- then do sanity check
- if you lost track for several consecutive frames, start searching from scratch again

• Drawing (and unwarping) the lane lines:

```
warp_zero = np.zeros_like(warped).astype(np.uint8)  
color_warp = np.dstack((warp_zero, warp_zero, warp_zero))
```

```
# Recast x and y into usable format for cv2.fillPoly()  
# The point must be in clockwise or counter clockwise (np.flipud())
```

```
pts_left = np.array([np.transpose(np.vstack([left_fitx, left_fity]))])
```

```
pts_right = np.array([np.flipud(np.transpose(np.vstack([right_fitx, right_fity])))])
```

```
pts = np.hstack((pts_left, pts_right))
```

```
cv2.fillPoly(color_warp, np.int_(pts), (0, 255, 0))
```

```
unwarp = cv2.warpPerspective(color_warp, Minv, (img.shape[1],  
img.shape[0]))
```

```
result = cv2.addWeighted(undist, 1, unwarp, 0.3, 0)
```

```
plt.imshow(result)
```



Steps

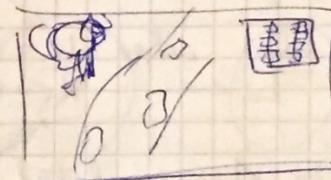
- ✓ Camera Calibration
- ✓ Distortion Correction
- ✓ Gradient, color, etc. thresholding
- ✓ perspective transform to **rectify** image
- ✓ Detect lane pixels and fit a polynomial to find lane boundaries
- ✓ Determine Lane curvature and vehicle offset from center
- ✓ Warp detected lane boundaries back onto the original image
- ✓ output lane boundaries and estimate of lane curvature and vehicle offset.

• Todo :

- search ahead
- line class
- check for too sharp curves
- sanity check
- smoothing

- adaptive thresholding

- Use Jupyter's interactive widgets for parameter tuning :)
- ↗ ipywidgets
- Add binary mini map to the



video to visualize your fitted lines and sliding window

L10: Neural Networks

Luis Soriano, Mat Leonard

1. April
22:00 Uhr

- Neural Networks, TensorFlow, DNNs, CNNs
- Linear regression \rightarrow line
- Linear regression is good for estimating values on a continuous spectrum
- Logistic Regression \rightarrow predicting discrete data
- linear reg. \rightarrow logistic reg. \rightarrow NNs.

• linear decision boundary (a line)

$$w_1x_1 + w_2x_2 + b = 0 \quad \text{weights}$$

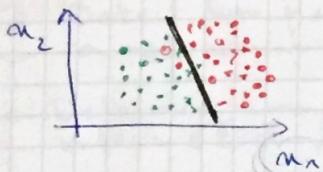
or

$$W \cdot X + b = 0 \quad \text{where } W = (w_1, w_2) \quad X = (x_1, x_2)$$

Bias

$y = 0$ or $1 \rightsquigarrow$ label

$$y = \begin{cases} 1 & \text{if } W \cdot X + b \geq 0 \\ 0 & \text{if } W \cdot X + b < 0 \end{cases} \quad \text{prediction}$$



- 3D data \Rightarrow decision boundary would be a plane

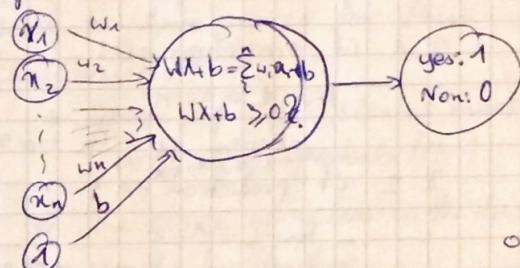
$$w_1x_1 + w_2x_2 + w_3x_3 + b = 0 \rightsquigarrow W \cdot X + b$$

$$W = (w_1, w_2, w_3)$$

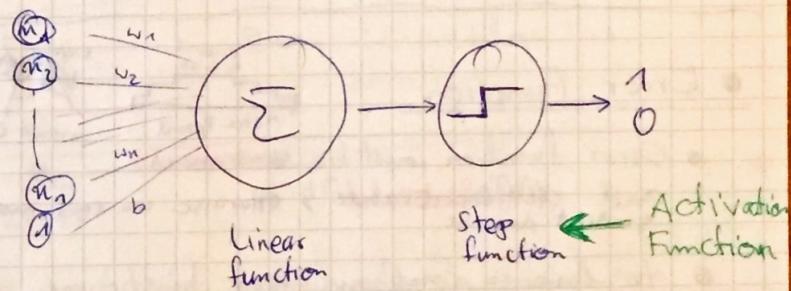
$$X = (x_1, x_2, x_3)$$

- N-dimensional data: x_1, \dots, x_n
↳ Boundary: n-1 dimensional hyperplane

② Perceptron:



- Step function: $y = f(n) = \begin{cases} 1 & \text{if } n \geq 0 \\ 0 & \text{if } n < 0 \end{cases}$



• Perceptron formula

$$f(x_1, \dots, x_n) = \begin{cases} 0 & \text{if } \sum w_i x_i + b < 0 \\ 1 & \text{if } \sum w_i x_i + b \geq 0 \end{cases}$$

③ Learning rate



Perception Algorithm

1. Start with random weights: w_1, \dots, w_n & b
 2. For every misclassified point (x_1, \dots, x_n)
 - 2.1 if prediction = 0 :
 - for $i = 1 \dots n$
 - change $w_i = w_i + \alpha x_i$
 - change $b = b + \alpha$
 - 2.2 if prediction = 1 :
 - for $i = 1 \dots n$
 - change $w_i = w_i - \alpha x_i$
 - change $b = b - \alpha$

repeat
till error
small
enough
or
for e.g.
100
times

① Error Function

- Error function must be continuous and differentiable; otherwise we can't use gradient descent
 - to have a continuous error function we need our predictions to be continuous and not discrete

④ Sigmoid function \rightarrow probabilities

• Multiclass classification

↳ softmax function

- When we have more than 2 classes, we need to have scores for them with following characteristics:

- sum of the scores must be 1
(because of probability)

- The score of class with higher probability must be higher

- let's divide the score by sum of all

Duck	Birres	Wahres
2	1	0
<u>2</u>	<u>1</u>	<u>0</u>
<u><u>2</u></u>	<u><u>1</u></u>	<u><u>0</u></u>
2+1=0	2+1=0	2+1=0

- good idea, but it has a problem:
What if some scores are negative
and we end up with sum being zero?

- ④ Let's take $\exp(x) \rightsquigarrow$ it is positive for every number

$$\exp(\mathbf{z}) > 0$$

$$\frac{e^2}{e^2 + e^1 + e^0}$$

$$\frac{e^x}{e^x + e^0}$$

$$\frac{e^0}{e^2 + e^1 + e^0}$$

• Softmax function:

Let's say we have n -classes and a linear ~~function~~ model that gives us the following scores z_1, \dots, z_n

$$P(\text{class } i) = \frac{e^{z_i}}{e^{z_1} + e^{z_2} + \dots + e^{z_n}}$$

• Soft max function for $n=2$ is the sigmoid function.

```
def softmax(L):    # L a list of numbers
    expL = np.exp(L)
    return np.divide(expL, expL.sum())
```

• One-hot encoding

• Maximum Likelihood

↳ we pick the model that gives the existing labels the highest probability.

⇒ By maximizing the probability, we can pick the best possible model.

- In order to tell how good or bad a model is, we calculate how likely it is. (its likelihood):
 - we calculate the probability of each point being the class it is according to the model.
 - multiply all these probabilities to obtain the probability, (the likelihood) of whole assignment

- to find the best model, we need to maximize the $P(\text{all})$ probability. $P(\text{all}) = P_1 \cdot P_2 \cdot P_3 \cdots P_n$ for n data points.

- But products are not good!! Let's take its log to turn product to sum

$$\Rightarrow P(\text{all}) = P_1 \cdot P_2 \cdot P_3 \cdots P_n$$

$$\Rightarrow \log(P(\text{all})) = \log(P_1 \cdot P_2 \cdots P_n) = \sum_i \log(P_i)$$

- We multiply the logs with -1, because $\log(x)$ is negative for $0 < x < 1$ (probabilities)

• This is called Cross Entropy

Sum of negatives of all probabilities of a model

$$-\sum_i \log(P_i)$$

- A good model has a smaller cross entropy!

- $-\log(P_i)$ for a point i is also a measure of error: the points that are misclassified have a larger value for $-\log(P_i)$

- now our goal is to minimize the cross entropy which is equivalent to maximizing the likelihood of the model

Events	Probabilities
unlikely?	How likely is it that the events happen according to the probabilities?
⇒ cross entropy high	likely? Cross entropy low

• Binary Cross Entropy Loss:

$$CE = - \sum_{i=1}^m y_i \ln(p_i) + (1-y_i) \ln(1-p_i)$$

$y_i = 1$ if class A
0 if not class A

• Multi-class Cross Entropy

$$CE = - \sum_{i=1}^n \sum_{j=1}^m y_{ij} \ln(p_{ij})$$

$y_{ij} = 0$ or 1

n = number of classes

• Logistic Regression.

- one of the most popular and useful algorithms in ML
- Building block of all that constitutes DL

• Logistic Regression:

- Take your data
- Pick a random model
- Calculate the error
- minimize the error, and obtain a better model
- repeat

$$\text{error} = -(1-y) \ln(1-\hat{y}) - y \ln(\hat{y})$$

$$\text{error function} = -\frac{1}{m} \sum_{i=1}^m [(1-y_i) \ln(1-\hat{y}_i) + y_i \ln(\hat{y}_i)]$$

event

$$\bullet \text{Multi-Class Entropy Error} : -\frac{1}{m} \sum_{i=1}^m \sum_{j=1}^n y_{ij} \ln(\hat{y}_{ij})$$

• Gradient Descent

$$w = \text{random()}$$

$$b = \text{random}$$

$$\hat{y} = \sigma(wx+b)$$

$$\nabla E = \left(\frac{\partial E}{\partial w}, \frac{\partial E}{\partial b} \right)$$

$$\alpha = 0.1 \text{ (learning rate)}$$

$$w' = w - \alpha \frac{\partial E}{\partial w}$$

$$b' = b - \alpha \frac{\partial E}{\partial b}$$

• Derivative of sigmoid $\sigma'(n) = \sigma(n)(1-\sigma(n))$

$$(e^n)' = e^n$$

$$\left(\frac{1}{n} \right)' = -\frac{1}{n^2}$$

$$\sigma(n) = \frac{1}{1+e^{-n}}$$

$$\frac{d}{dn} \left(\frac{1}{1+e^{-n}} \right) = \frac{e^{-n}}{(1+e^{-n})^2} = \frac{1}{1+e^{-n}} \cdot \frac{e^{-n}}{1+e^{-n}} = \sigma(n) \cdot (1-\sigma(n))$$

• We are after $\frac{\partial E}{\partial w}$ and $\frac{\partial E}{\partial b}$

- for m data points $x^{(1)}, x^{(2)}, \dots, x^{(m)}$ the error is $E = -\frac{1}{m} \sum_{i=1}^m (y_i \ln(\hat{y}_i) + (1-y_i) \ln(1-\hat{y}_i))$

We want to calculate gradient of E at point $n = (n_1, n_2, \dots, n_n)$ given by the partial derivative: \rightarrow

$$\leftarrow \nabla E \left(\frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n}, \frac{\partial E}{\partial b} \right)$$

- Let's actually think of the error that each point generates and calculate the derivative of this error. The total error is then the average of errors at all points. The error produced by each point is

$$E = -y \ln(\hat{y}) - (1-y) \ln(1-\hat{y})$$

~~Calculate~~

- In order to calculate the derivative of this error w.r.t. the weights, let's first calculate $\frac{\partial \hat{y}}{\partial w_j}$

$$\hat{y} = \sigma(wx + b)$$

$$\begin{aligned} \frac{\partial}{\partial w_j} (\hat{y}) &= \frac{\partial}{\partial w_j} \sigma(wx+b) \\ &= \sigma'(wx+b)(1-\sigma(wx+b)) \cdot \frac{\partial}{\partial w_j} (wx+b) \\ &= \hat{y}(1-\hat{y}) \cdot \frac{\partial}{\partial w_j} (wx+b) \\ &= \hat{y}(1-\hat{y}) \cdot \frac{\partial}{\partial w_j} (w_1x_1 + \dots + w_jx_j + \dots + w_nx_n + b) \\ &= \hat{y}(1-\hat{y}) \cdot x_j \end{aligned}$$

because the only term in the sum which is not constant w.r.t. w_j is $w_j x_j$, which has derivative x_j

- Now we can go ahead and calculate the derivative of the error E at point x w.r.t. w_j



$$\begin{aligned} \frac{\partial E}{\partial w_j} &= \frac{\partial}{\partial w_j} [-y \ln(\hat{y}) - (1-y) \ln(1-\hat{y})] \\ &= -y \frac{\partial}{\partial w_j} \ln(\hat{y}) - (1-y) \frac{\partial}{\partial w_j} \ln(1-\hat{y}) \\ &= -y \cdot \frac{1}{\hat{y}} \frac{\partial}{\partial w_j} \hat{y} - (1-y) \cdot \frac{1}{1-\hat{y}} \cdot \frac{\partial}{\partial w_j} (1-\hat{y}) \\ &= y \cdot \frac{1}{\hat{y}} \hat{y}(1-\hat{y}) x_j - (1-y) \cdot \frac{1}{1-\hat{y}} \cdot (-1)\hat{y}(1-\hat{y}) x_j \\ &= -y(1-\hat{y}) \cdot x_j + (1-y)\hat{y} x_j \\ &= -(y-\hat{y}) x_j \end{aligned}$$

Similarly

$$\frac{\partial E}{\partial b} = -(y-\hat{y})$$

- This has a very important implication: For a point with coordinates (x_1, \dots, x_n) , label y , and prediction \hat{y} the gradient of the error function at that point is $-(y-\hat{y})(x_1, \dots, x_n, 1)$

$$\Rightarrow \nabla E = -(y-\hat{y})(x_1, \dots, x_n, 1)$$

- The gradient is actually a scalar times the coordinates of the point. And this scalar is nothing more than a difference multiple of the between the label and the prediction.

\Rightarrow The closer the prediction to label, the smaller the gradient

or

The farther the label from prediction, the larger the gradient

- So exactly like the perceptron algorithm, small gradient means change the coordinates a little bit, and larger gradient means change the coordinates a lot.

Gradient Descent Step

$$w'_i \leftarrow w_i - \alpha [-(y - \hat{y})x_i]$$

$$\Rightarrow w'_i \leftarrow w_i + \alpha (y - \hat{y})x_i$$

$\Delta w_i = \alpha \delta x_i$
δ error term

and $b' \leftarrow b + \alpha (y - \hat{y})$

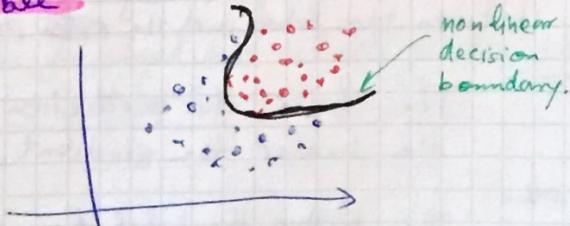
Grad. Descent vs. Perceptron Alg.

Change
 $w_i \rightarrow w_i + \alpha (y - \hat{y})x_i$ changing w_i to $\begin{cases} w_i + \alpha x_i & \text{label positive} \\ w_i - \alpha x_i & \text{label negative} \end{cases}$

- Neural Networks can show their full potential with non-linear data

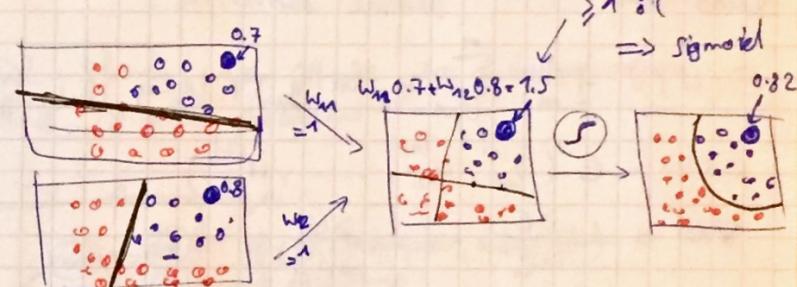
- **Linear data** : when a line can separate the data

- **Non-linearly separable**

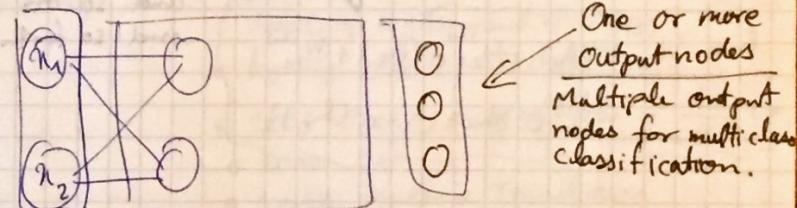


- Lovely explanation :)

We want to build (linear) combinations of several models to achieve new ones that separate complex data



- we can add weights to each model too to make its contribution more important.



input layer hidden layer output layer ← where linear models get combined to obtain a non-linear model

multiple output

- softmax to get probability score of each class.
- Feed forward

- Backprop
- Chain rule: derivative of composition can be shown as multiplying the partial derivatives.

$$E(W) = -\frac{1}{m} \sum_{i=1}^m y_i \ln(\hat{y}_i) + (1-y_i) \ln(1-\hat{y}_i)$$

$$E(W) = E(W^{(1)}, W^{(2)}, \dots, W^{(L)})$$

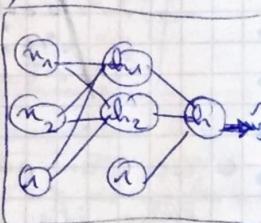
$$\nabla E = \left(\frac{\partial E}{\partial W_1^{(1)}}, \dots, \frac{\partial E}{\partial W_L^{(L)}} \right)$$

$$\frac{\partial E}{\partial W_1^{(1)}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial h} \frac{\partial h}{\partial h_1} \frac{\partial h_1}{\partial W_1^{(1)}}$$

$$\frac{\partial E}{\partial \hat{y}} = (\hat{y} - y)$$

$h = W^{(2)}_M \sigma(h_1) + W^{(2)}_{21} \sigma(h_2) + W^{(2)}_{31}$

$$\Rightarrow \frac{\partial h}{\partial h_1} = W^{(2)}_{11} \sigma'(h_1)(1 - \sigma(h_1))$$



and so on
and so forth

LM : TensorFlow

09. April 22:00
2h

- Vincent Vanhoucke: principle scientist at Google's DL and director of robotics.
- DL shines when there are a lot of data and complex problems
- DL Advantages:
 - a lot of techniques, adaptable to many different problems.
 - common infrastructure
 - and a - common language to describe things

- History
 - LeCun's LeNet end of 90's detecting digits handwriting
 - beginning 2000's: Alex Net
 - 2009 speech recognition
 - 2012 CV
 - 2014 Machine Translation

what changed? [at end of 2000's NNs made a comeback]

- a lot of data
- fast GPUs

- We will use TF to classify images of **not MNIST** dataset: ~~a~~ a dataset of images of English letters A-J

```

import tensorflow as tf
hello_const = tf.constant("hello")
with tf.Session() as sess:
    output = sess.run(hello_const)
    print(output)

```

- TF data is not stored as integers, floats, or strings.
The values are encapsulated in an object called
Tensors

- hello_const \rightsquigarrow 0-dimensional string tensor.
 \rightsquigarrow a constant tensor
- $A = \text{tf. constant}(1234)$ \rightsquigarrow 0-dim int32 tensor
- $B = \text{tf. constant}([123, 456, 789])$ \rightsquigarrow 1D int32 tensor
- $C = \text{tf. constant}(([123, 456, 789], [2, 3, 4]))$ \rightsquigarrow 2D \approx \approx

~~Variables~~

Computational Graph

- `tf.Session()` an environment for running computational graphs

tf.placeholder(), feed-dict

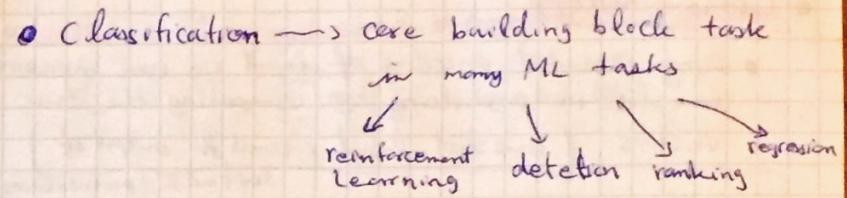
```
x = tf.placeholder(tf.string)
```

With `tf.Session()` as sess:

```
output = sess.run(x, feed_dict={x: "hello"})
```

tf.add(), tf.subtract(), tf.multiply()

tf.cast() \rightsquigarrow no implicit casting!



Logistic Classifiers (Linear Classifier)

$$Wx + b = \hat{y}$$

$$\text{Softmax function: } S(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$

y_i, y_j : scores

Softmax function takes the **scores** (a.k.a. **logits**) and turns them **to probabilities**.

- Tensorflow uses $xW + b$ instead of $Wx + b$

- `tf.Variable()`, $a = \text{tf.Variable}(5)$

~~tf.global_variables_initializer()~~

- `tf.global_variables_initializer()`

- `tf.truncated_normal()` \rightsquigarrow generate random numbers from Normal dist
($n_features, n_timesteps$)

- `tf.zeros()`

- `tf.matmul()`

- `tf.nn.softmax()`